RESEARCH PREPRINT

Fast Evaluation of Truncated Neumann Series by Low-Product Radix Kernels

Piyush Sao

Oak Ridge National Laboratory, Oak Ridge, TN, USA

Abstract

Truncated Neumann series $S_k(A) = I+A+\cdots+A^{k-1}$ are a standard route to approximate inverses and polynomial preconditioners, but high-degree truncations remain expensive in dense linear algebra because matrix-matrix products dominate the cost. Classical splitting methods reduce the work to $O(\log k)$, so further progress depends on finding better short product formulas for the geometric block $I + B + \cdots + B^{m-1}$. Beyond the known radix-5 case, explicit higher-radix kernels were unavailable, and approximate kernels lacked a rigorous iterative framework.

We make three contributions. First, we derive an exact radix-9 kernel that evaluates $I + B + \cdots + B^8$ in three matrix products, giving the first exact construction beyond radix 5. Second, for radices 15 and 24 we obtain certified approximate kernels whose low-degree coefficients are exact and whose spillover is rigorously bounded. Third, we develop a residual-based radix-kernel framework that shows how to use these approximate kernels repeatedly without losing control of the error: it gives explicit conditions under which the iteration converges and the residual decays, including for nonnormal matrices.

Under the dense-matrix GEMM-count model, the new kernels reduce the asymptotic matrix-product count by roughly one quarter relative to repeated squaring. Theory-driven experiments validate the predicted stability thresholds, illustrate the behavior on normal and nonnormal examples, and confirm the expected reduction in matrix products.

Keywords: Neumann series; matrix polynomial; radix kernel; matrix multiplication; computational complexity

1. Introduction

The truncated Neumann series is a core object in numerical linear algebra. For a matrix A with spectral radius $\rho(A) < 1$, it is defined by

$$S_k(A) = \sum_{j=0}^{k-1} A^j = I + A + A^2 + \cdots + A^{k-1}. \tag{1}$$

As $k \to \infty$, $S_k(A)$ converges to $(I - A)^{-1}$, so finite truncations give practical approximate inverses. These approximations appear in polynomial preconditioning [1, 2], resolvent approximation [3], log-determinant estimation [4], and massive MIMO signal detection [5]. In dense problems, the dominant cost is matrix-matrix multiplication (GEMM), so reducing product count is the main efficiency goal.

Existing splitting methods already reduce the naive $k-1$ products to $O(\log k)$ [6]. Binary splitting (repeated squaring) uses $2\log_2 k$ products. Hybrid radix-$\{2, 3\}$ schemes [7] lower the constant only slightly, to about $1.9\log_2 k$. Thus asymptotic order is no longer the bottleneck; the leading coefficient is.

We therefore adopt dense-matrix GEMM count as a surrogate complexity model. The paper optimizes matrix-product counts rather than wall-clock time, and the numerical section is designed to validate the resulting coefficient, product-count, and stability consequences of the constructions rather than to serve as an application-scale performance study.

Improving that coefficient therefore requires better radix kernels. Let $T_m(B) = I + B + \cdots + B^{m-1}$ and let μ_m denote the minimum number of products to evaluate T_m. The decomposition

$$S_{mn}(A) = S_n(A) \cdot T_m(A^n). \tag{2}$$

shows that one radix-m update costs $C(m) = \mu_m + 2$ products: μ_m for the kernel, one for concatenation, and one for updating the next power. Hence the asymptotic coefficient is $C(m)/\log_2 m$. Each product can at most double polynomial degree, so $\mu_m \geq \lceil \log_2(m-1) \rceil$. This lower bound gives $\mu_9 \geq 3$ and $\mu_{15} \geq 4$, which suggests that

This manuscript has been authored by UT-Battelle, LLC under Contract No. DE-AC05-00OR22725 with the U.S. Department of Energy. The publisher, by accepting the article for publication, acknowledges that the United States Government retains a non-exclusive, paid-up, irrevocable, world-wide license to publish or reproduce the published form of this manuscript, or allow others to do so, for United States Government purposes. The Department of Energy will provide public access to these results of federally sponsored research in accordance with the DOE Public Access Plan (http://energy.gov/downloads/doe-public-access-plan).

higher-radix kernels could improve constants if such kernels exist. Gustafsson et al. [5] achieved $\mu_5 = 2$ for radix 5, but explicit constructions beyond radix 5 were missing. This leaves two linked questions: can radix 9 reach its lower bound exactly, and can higher-radix kernels remain useful when they are only approximate?

We answer these questions with four contributions. We construct an exact radix-9 kernel with $\mu_9 = 3$ and rational coefficients (Theorem 4.1), giving $5\log_9 k = 1.58\log_2 k$ products. This is the first exact construction beyond radix 5 and gives a 21% improvement over binary splitting. For radices 15 and 24, we combine numerical discovery with interval certification to obtain prefix-exact higher-radix kernels with controlled spillover. We develop a residual-based general radix-kernel framework that accommodates spillover while preserving coefficient $(\mu_m + 2)/\log_2 m$. We prove matrix-level stability bounds for that framework, including normal, numerical-range, and pseudospectral corollaries, and validate them numerically. Within this framework, radix 15 reaches $6/\log_2 15 \approx 1.54$ and radix 24 reaches $7/\log_2 24 \approx 1.53$, the smallest asymptotic coefficients achieved by the constructions in this paper.

The paper proceeds as follows. Section 2 reviews background and the series-specific prior work. Section 3 places the kernel viewpoint in the broader polynomial-evaluation literature. Section 4 presents the radix-9 construction. Section 5 separates numerical search from certified approximate higher-radix kernels. Section 6 develops the general framework and its cost analysis. Section 7 develops matrix-level stability theorems, and Section 8 validates those theorems numerically while comparing product counts.

2. Background

We review existing methods and show that naive kernel computation is their main limitation. Given $A \in \mathbb{R}^{d\times d}$ with $\rho(A) < 1$, define the truncated Neumann series as in (1):

$$S_k(A) = \sum_{j=0}^{k-1} A^j = I + A + A^2 + \cdots + A^{k-1}. \tag{3}$$

We measure cost by counting matrix products (GEMMs), which dominate when $d \gtrsim 100$ for dense matrices.

2.1 Splitting Methods

Several radix-based schemes exist, differing in how many terms they add per step.

- **Naive evaluation.** The naive approach accumulates powers $I, A, A^2, \ldots$, requiring $k-1$ matrix products—prohibitive for large k.
- **Radix-m splitting.** Faster algorithms use the factorization $S_{mn}(A) = S_n(A)\cdot(I + A^n + \cdots + A^{(m-1)n})$. A radix-$m$ scheme iterates this identity. Each step computes:
  - powers $A^{2n}, \ldots, A^{(m-1)n}$ from A^n ($m-2$ products),
  - the product $S_n \cdot (I + A^n + \cdots + A^{(m-1)n})$ (1 product),
  - and $A^{mn} = A^{(m-1)n}\cdot A^n$ (1 product).

  The naive cost per step is $C(m) = m$ products, giving coefficient $m/\log_2 m$.
- **Binary splitting.** For $m = 2$, the identity simplifies to $S_{2n} = S_n(I + A^n)$, using only 2 products per step [8, 6]: (i) $A^{2n} = A^n \cdot A^n$; (ii) $S_{2n} = S_n \cdot (I + A^n)$. Total: $2\log_2 k$ products.
- **Ternary splitting.** For $m = 3$, $S_{3n} = S_n(I + A^n + A^{2n})$: (i) $A^{2n} = A^n \cdot A^n$; (ii) $A^{3n} = A^{2n}\cdot A^n$; (iii) $S_{3n} = S_n \cdot (I + A^n + A^{2n})$. Total: $3\log_3 k \approx 1.89\log_2 k$ products.

Optimal radix for naive splitting. The coefficient $C(m)/\log_2 m$ equals $m/\log_2 m$, minimized at $m = e \approx 2.718$; thus $m = 3$ is optimal among integers:

Radix m	Cost $C(m)$	Coefficient $C(m)/\log_2 m$
2	2	$2/\log_2 2 = 2.00$
3	3	$3/\log_2 3 \approx 1.89$
4	4	$4/\log_2 4 = 2.00$
5	5	$5/\log_2 5 \approx 2.15$
9	9	$9/\log_2 9 \approx 2.84$
15	15	$15/\log_2 15 \approx 3.84$
24	24	$24/\log_2 24 \approx 5.23$

Thus ternary has the best coefficient among pure radix methods. Dimitrov and Cooklev [7] extended this to hybrid radix-$\{2,3\}$ methods for arbitrary k, achieving coefficient ≈ 1.9. These mixed $\{2,3\}$ strategies are near-optimal, but only with naive kernel computation.

2.2 Radix Kernels

The splitting methods above compute the partial sum $T_m(B) = I + B + \cdots + B^{m-1}$ naively, requiring $m-2$ products for powers $B^2, \ldots, B^{m-1}$. We call T_m the *kernel*. If the kernel cost μ_m grew only logarithmically in m, the coefficient $(\mu_m + 2)/\log_2 m$ could fall below 1.89. Two identities make this possible:

- **Block concatenation.** $S_{mn}(A) = S_n(A)\cdot T_m(A^n)$ for any $m, n \geq 1$, verified by expanding both sums.
- **Telescoping.**

$$T_m(B)(I - B) = I - B^m \tag{4}$$

  This geometric-series identity extracts B^m from the kernel without extra products.

Cost model. Let μ_m denote the minimum number of products to compute $T_m(B)$. A radix-m update $(S_n, A^n) \to (S_{mn}, A^{mn})$ uses:

- μ_m products for the kernel,
- 1 product for telescoping: $I - A^{mn} = T_m(A^n)(I - A^n)$,
- 1 product for concatenation: $S_{mn} = S_n \cdot T_m(A^n)$.

The total is $C(m) = \mu_m + 2$, giving coefficient $C(m)/\log_2 m$. Later, Sections 6 and 6.4 replace the exact telescoping/power-update product with residual recomputation for approximate kernels; the bookkeeping changes, but the total remains $\mu_m + 2$.

Lower bound. Since each product at most doubles polynomial degree, p products starting from $\{I, B\}$ reach degree at most 2^p. Because T_m has degree $m - 1$, evaluating it requires at least $\lceil \log_2(m-1) \rceil$ products. In particular, $\mu_5 \geq 2$ and $\mu_9 \geq 3$. Can these bounds be achieved?

2.3 Prior Work: Quinary (Radix-5) Kernel

For radix-5, the answer is yes. Gustafsson et al. [5] achieved $\mu_5 = 2$ with the following construction: (i) $U := B^2$; (ii) $V := U(B + U) = B^3 + B^4$; (iii) $T_5(B) := I + B + U + V$. Direct verification confirms $T_5(B) = I + B + B^2 + B^3 + B^4$.

With $\mu_5 = 2$, a quinary update costs $C(5) = 4$ products, giving coefficient $4/\log_2 5 \approx 1.72$.

Open questions. The quinary kernel reduces the coefficient, but gaps remain. First, no radix-9 construction exists, though the degree-doubling bound permits $\mu_9 = 3$, giving coefficient $5/\log_2 9 \approx 1.58$. Second, it is unknown whether exact rational kernels exist for higher radices such as $m = 15$. Approximate constructions may have nonzero *spillover* into degrees $\geq m$, breaking the telescoping identity (4) and preventing exact extraction of B^m. No framework exists for such kernels in series algorithms. These gaps set the agenda for the rest of the paper. Section 3 places them in the broader matrix-polynomial-evaluation literature. Sections 4–6 then develop exact and approximate higher-radix kernels, and the later sections analyze their stability and numerical performance.

3. Related Work

The background above covered splitting methods specific to truncated Neumann series. A second line of context comes from general matrix-polynomial evaluation, where special structure can justify product counts that generic schemes cannot reach.

Within Neumann-series evaluation itself, Westreich [8] obtained $3\log_2 k$ products by binary splitting. Lei and Nakamura [6] reduced this to $2\log_2 k - 1$ by reusing intermediate computations. Dimitrov and Cooklev [7] then introduced hybrid radix-$\{2, 3\}$ methods and, under naive kernel evaluation, argued that higher radices offered no practical improvement. Matsumoto, Takagi, and Takagi [9] later showed that higher radices can improve the leading coefficient, but they did not provide explicit kernel circuits or a framework for deciding which radices are useful.

For general matrix polynomials, the Paterson–Stockmeyer (PS) method [10] achieves $O(\sqrt{n})$ products for degree-n evaluation, and Fasi and Higham [11] proved optimality within the PS scheme. Composition methods can do better on special polynomial families: Sastre and Ibáñez [12] achieved degree 15 in 4 products, exceeding the degree 9 reachable by PS with 4 products. That comparison is not directly transferable to Neumann kernels. The radix-15 Neumann kernel $T_{15}(B) = I + B + \cdots + B^{14}$ has degree 14 and fixed unit coefficients, whereas the Sastre–Ibáñez constructions exploit the more flexible coefficient structure of exponential-family polynomials. For Neumann kernels, the uniform-coefficient constraint leaves much less freedom in the factorization, so a degree-15 result for a different polynomial family does not imply a 4-product construction for T_{15}.

Generic schemes such as Horner evaluation or baby-step/giant-step reorganize arbitrary polynomial evaluation, but they do not exploit the fixed unit-coefficient geometric structure of $T_m(B)$. For that reason they provide useful baselines for general matrix polynomials, yet they do not by themselves explain the low-product kernel circuits sought here. Likewise, Padé-type approximate inverses belong to a rational-approximation framework with different approximation classes and solve costs, so they lie outside the polynomial-kernel scope of this paper.

Our radix kernels import that composition viewpoint into the geometric-series setting. In particular, they extend Gustafsson et al.'s quinary construction [5] to higher radices, and the accompanying kernel-to-update framework shows how better kernel circuits translate directly into faster Neumann-series algorithms.

4. Radix-9 Kernel

The quinary kernel [5] achieves coefficient $4/\log_2 5 \approx 1.72$. The next natural target is radix 9. To improve on 1.72, a radix-9 update must use at most 5 products. Since $C(9) = \mu_9 + 2$, this means $\mu_9 \leq 3$. The degree-doubling bound already gives $\mu_9 \geq 3$, so a 3-product kernel would be optimal.

Derivation strategy. To hit this bound, we build the third product around $U = B^2$ and $V = U(B + 2U) = B^3 + 2B^4$, which are available after two products. We then seek one additional product $P \cdot Q$, with P and Q linear combinations of $\{B, U, V\}$, whose convolution fixes the high-degree terms $B^5, \ldots, B^8$. The remaining degrees B^2, B^3, B^4 can then be corrected using only additions of U and V. Matching the target coefficients gives an 8×8 linear system for the coefficients of P and Q.

Theorem 4.1 (Radix-9 kernel in 3 products)**.** *The kernel $T_9(B) = I + B + \cdots + B^8$ is computed by:*

1. $U := B^2$ *(MM1)*
2. $V := U(B + 2U) = B^3 + 2B^4$ *(MM2)*
3. $P := \frac{3}{20}B + 2U + V, \quad Q := \frac{11}{40}B - \frac{1}{8}U + \frac{1}{4}V$
4. $W := P \cdot Q$ *(MM3)*
5. $T_9(B) := W + I + B + \frac{767}{800}U + \frac{15}{32}V$

One carefully chosen product $P \cdot Q$ simultaneously generates four high-degree terms (B^5 through B^8), while the low-degree corrections adjust B^2–B^4 using only additions.

Proof. Let $\mathbf{p} = [\frac{3}{20}, 2, 1, 2]$ and $\mathbf{q} = [\frac{11}{40}, -\frac{1}{8}, \frac{1}{4}, \frac{1}{2}]$ be the coefficients of P and Q in powers B^1, B^2, B^3, B^4. Expanding $P \cdot Q$ shows that the coefficients of B^5 through B^8 are exactly 1. The same expansion gives coefficients $\frac{33}{800}$, $\frac{17}{32}$, and $\frac{1}{16}$ on B^2, B^3, and B^4. Adding $\frac{767}{800}U + \frac{15}{32}V$ raises those coefficients to 1, so the resulting polynomial is $I + B + \cdots + B^8$. □

Remark 4.2 (Rational coefficients). The radix-9 kernel uses only rational coefficients (fractions with denominators dividing 800), enabling exact computation in rational or fixed-point arithmetic. This contrasts with radix-15 (Section 5), where numerical optimization yields only real coefficients with nonzero "spillover" into higher-degree terms. The transition from exact rational solutions (radix-5, radix-9) to approximate real solutions (radix-15) suggests that any exact higher-radix construction, if it exists, is substantially harder to uncover.

With $\mu_9 = 3$, the update costs $C(9) = 5$ products. Iterating $t = \log_9 k$ times gives $S_k(A)$ in $5 \log_9 k \approx 1.58 \log_2 k$ products—a 21% improvement over binary and 8% over quinary.

In summary, the radix-9 kernel achieves the smallest asymptotic coefficient among the exact constructions discussed here, at the cost of one additional matrix product per update. The next section explores whether higher radices can improve further.

5. Certified Approximate Higher-Radix Kernels

The radix-9 kernel is *exact* and lower-bound optimal. For higher radices, the same degree-counting argument still points to useful coefficients—$\mu_{15} \geq 4$ and $\mu_{24} \geq 5$—but it does not tell us whether exact kernels exist or how approximate ones should be used. This section therefore separates two logically distinct tasks: *numerical discovery* of promising circuits, and *certified existence* of nearby prefix-exact kernels. Keeping those tasks separate prevents an optimizer output from being mistaken for a theorem.

5.1 Search Ansatz

We use the bilinear chained-product ansatz already implicit in the exact radix-9 construction. Fix $P_1 = B^2$. For a K-product kernel, define basis sets

$$\mathcal{B}_1 = \{I, B, P_1\}, \qquad \mathcal{B}_{i+1} = \mathcal{B}_i \cup \{P_i\},$$

and for $i = 2, \ldots, K$ write

$$L_i = \sum_{X \in \mathcal{B}_i} \ell_{i,X} X, \qquad R_i = \sum_{X \in \mathcal{B}_i} r_{i,X} X, \qquad (5)$$
$$P_i = L_i R_i.$$

The final kernel is

$$f(B) = I + \gamma_1 B + \gamma_2 P_1 + \cdots + \gamma_K P_{K-1} + P_K. \qquad (6)$$

For radix 15 we use $K = 4$, giving 28 parameters. For radix 24 we use $K = 5$, giving 41 parameters.

5.2 Numerical Discovery

Let $[f(B)]_j$ denote the coefficient of B^j. The discovery objective is the prefix residual

$$\Phi_m(\theta) = \sum_{j=0}^{m-1} \left([f_\theta(B)]_j - 1\right)^2,$$

where θ collects all ansatz parameters. We minimize Φ_m by multi-start local optimization. This stage is deliberately heuristic: it identifies promising approximate solutions but does not by itself justify theorem-grade claims.

For radix 15 the best 4-product approximate solution drives the prefix error below 5×10^{-16} while leaving spillover at degrees 15 and 16. For radix 24 we refine the active prefix subsystem to a 16-decimal circuit before fixing the final approximate solution printed in the appendix.

5.3 Certified Approximate Kernels

The next step resolves the main expectation created by the search: each promising approximate solution is converted into a *certified* kernel by solving the prefix equations exactly in a nearby parameter box. Write

$$F_m(\theta) = ([f_\theta]_0 - 1, \ldots, [f_\theta]_{m-1} - 1).$$

At each approximate solution θ_0 we compute $DF_m(\theta_0)$, select a well-conditioned square subsystem, and apply a Krawczyk inclusion test on a small nearby box in those variables while holding the remaining variables fixed.

Theorem 5.1 (Computer-assisted certified radix-15 and radix-24 kernels)**.** *There exists a certified radix-15 kernel in the 4-product bilinear ansatz and a certified radix-24 kernel in the 5-product bilinear ansatz such that the prefix equations $F_{15}(\theta) = 0$ and $F_{24}(\theta) = 0$ hold exactly. The corresponding verified parameter boxes have maximum coordinate radii below 4.5×10^{-14} and 6.2×10^{-14}, respectively. Evaluating the kernels and their error maps over those boxes yields interval enclosures for every displayed spillover coefficient.*

Proof. This is a computer-assisted proof. For each radix, the displayed approximate solution determines a square subsystem of the prefix equations whose Krawczyk operator is verified to map a nearby parameter box strictly into itself. The Krawczyk inclusion theorem therefore yields an exact root of the full prefix system in that box. Interval evaluation of the corresponding kernel and error map over the same box then gives the stated spillover enclosures. □

The certified Krawczyk subsystem has 15 active variables for radix 15 and 24 active variables for radix 24. The corresponding condition estimates are 3.83×10^1 and 2.38×10^2, and the certified maximum box radii are 4.5×10^{-14} and 6.2×10^{-14}. The full active sets, interval coefficients, and radii are recorded in the supplementary certification files `certified_coeffs_radix15.json` and `certified_coeffs_radix24.json`.

For radix 15, the certified error map is centered at

$$\begin{aligned} E_{15}(z) = {} & 0.81479100\, z^{15} \\ & - 0.27315103\, z^{16} \\ & + 0.45836003\, z^{17}, \end{aligned}$$

with interval radii below 2.1×10^{-13}. For radix 24, the certified error map is centered at

$$\begin{aligned} E_{24}(z) = {} & 0.35517021\, z^{24} + 0.43096456\, z^{25} \\ & + 0.16540498\, z^{26} - 0.13946877\, z^{27} \\ & - 0.24168001\, z^{28} - 0.10159049\, z^{29} \\ & + 0.10987210\, z^{30} + 0.20681982\, z^{31} \\ & + 0.15756340\, z^{32} + 0.05694420\, z^{33}, \end{aligned}$$

with interval radii below 2.8×10^{-12}. These certified midpoint coefficients are used in Sections 7 and 8; the radix-24 appendix table prints the corresponding 16-decimal circuit.

5.4 Local Structure

Certification also yields local geometric information about the solution set, clarifying why approximate kernels are natural in these ansätze. For the certified radix-15 kernel, the prefix Jacobian has full row rank 15 in the 28-parameter ansatz, so the local prefix-exact set is a smooth manifold of dimension $28 - 15 = 13$. For radix 24, the corresponding dimension is $41 - 24 = 17$. Appending the leading spillover coefficient as one more constraint raises the rank by one in both cases, so the local dimension drops by one.

This is still only a *local* result: it explains why prefix exactness plus small spillover is a natural optimization target near the certified solutions, but it does not prove global nonexistence of exact higher-radix kernels. The corresponding Jacobian-rank statement is recorded in Appendix 5.

The remainder of the paper treats these certified approximate kernels as mathematical objects rather than optimizer outputs. Section 6 shows how spillover is absorbed by residual iteration, and Section 7 turns the certified error maps into matrix-norm decay bounds.

6. General Radix-Kernel Framework

Among the exact kernels developed here, radix 9 has the smallest coefficient, $5/\log_2 9 \approx 1.58$. Higher radices can do better in principle. A 4-product radix-15 kernel would drop the coefficient to $6/\log_2 15 \approx 1.54$, and a 5-product radix-24 kernel to $7/\log_2 24 \approx 1.53$. Our certified higher-radix kernels reach those coefficients, but they do so with nonzero *spillover* at degrees ≥ 15, so the standard exact update no longer applies. The issue is therefore not whether higher-radix kernels help, but how to use approximate ones without giving up the asymptotic gain.

Spillover invalidates the telescoping identity $T_m(A^n)(I - A^n) = I - A^{mn}$ that normally lets us extract A^{mn}. With spillover the identity fails and the power cannot be recovered. Hensel lifting ($Y' = Y(2I - MY)$) sidesteps power extraction, but its coefficient is 2.00—no better than binary splitting. We need an iteration that keeps the 1.54–1.53 promise of higher-radix kernels without exact power extraction.

This section resolves both obstacles by applying the kernel to the *residual* rather than to powers of A, eliminating the need for power extraction. It proceeds in three

stages: formalize approximate kernels, analyze their composition as polynomials, then lift to matrices.

6.1 Approximate Radix Kernels

The following definition makes precise what "approximate kernel" means and how its error behaves. Although the kernels constructed here have real coefficients, the residual iteration and functional calculus are most naturally stated over $\mathbb{C}$.

Definition 6.1 (Approximate radix-m kernel)**.** A polynomial $f \in \mathbb{C}[z]$ is an *approximate radix-m kernel* if $(1-z)f(z) = 1+O(z^m)$. Equivalently, $f(z) \equiv (1-z)^{-1} \pmod{z^m}$. The terms of degree $\geq m$ in f are called *spillover*. When the spillover is zero, the kernel is *exact* and coincides with $T_m(z)$.

Definition 6.2 (Error map)**.** Given an approximate radix-m kernel f, define $E(z) = 1-(1-z)f(z)$.

Lemma 6.3 (Error order)**.** *If $f(z) = \sum a_j z^j$ is an approximate radix-m kernel and $c := 1 - a_m$, then*

$$E(z) = c\,z^m + O(z^{m+1}).$$

If f is exact, $a_m = 0$, so $c = 1$ and $E(z) = z^m$.

Proof. Since $(1-z)f(z) = 1 + O(z^m)$, the degree-m term is $(1-z)f(z) = 1 + (a_m - a_{m-1})\,z^m + \cdots = 1 + (a_m - 1)\,z^m + \cdots$, so $E(z) = 1-(1-z)f(z) = (1-a_m)\,z^m + O(z^{m+1})$. □

6.2 General Radix-Kernel Summation

Applying the kernel to its own error, $f(E(z))$, corrects the residual left by $f(z)$.

Definition 6.4 (General radix-kernel summation)**.** Given an approximate radix-m kernel f with error map E, one *lift* of the general radix-kernel summation produces

$$f(z)\cdot f\big(1-(1-z)f(z)\big) = f(z)\cdot f(E(z)). \tag{7}$$

Lemma 6.5 (Composition identity)**.** *For any approximate radix-m kernel f with error map E:*

$$E(E(z)) = 1-(1-z)f(z)f(E(z)). \tag{8}$$

Proof. By definition, $1-E(z) = (1-z)f(z)$. Applying the definition of E at the point $E(z)$:

$$E(E(z)) = 1-(1-E(z))f(E(z)) = 1-(1-z)f(z)f(E(z)).$$

□

Lemma 6.6 (Prefix growth)**.** *If f is an approximate radix-m kernel, then $f(z)\cdot f(E(z))$ matches the first m^2 coefficients of $(1-z)^{-1}$.*

This is the core composition mechanism behind the framework: one lift upgrades an m-term exact prefix to an m^2-term exact prefix, and repeated use of the lemma is what drives the m^n growth used later in the cost and stability analysis.

Proof. By Lemma 6.5, $(1-z)f(z)f(E(z)) = 1-E(E(z))$. By Lemma 6.3, $E(z) = O(z^m)$, so $E(E(z)) = O(z^{m^2})$. Hence $(1-z)f(z)f(E(z)) = 1 - O(z^{m^2})$, and $f(z)f(E(z))$ agrees with $(1-z)^{-1}$ to order m^2-1. □

Lemma 6.7 (Error propagation)**.** *Write $E^{[n]}$ for the n-fold composition ($E^{[0]}(z) = z$, $E^{[n+1]} = E\circ E^{[n]}$). Let $E(z) = c\,z^m + O(z^{m+1})$. Then*

$$E^{[n]}(z) = c^{(m^n-1)/(m-1)}\,z^{m^n} + O(z^{m^n+1}).$$

Spillover is harmless to the prefix: regardless of c, it is pushed to degree $\geq m^n$ after n steps.

Proof. $E(E(z)) = c\,(cz^m)^m + \cdots = c^{m+1}\,z^{m^2} + \cdots$. By induction, $E^{[n]}(z) = c^{1+m+\cdots+m^{n-1}}\,z^{m^n} + \cdots = c^{(m^n-1)/(m-1)}\,z^{m^n} + \cdots$. □

Remark 6.8 (Analytic stability). In floating-point arithmetic, the coefficient magnitude matters. When $|c| < 1$, the leading coefficient $c^{(m^n-1)/(m-1)}$ decays, bounding the residual norm. For the radix-15 kernel (Section 5), $[f]_{15} \approx 0.19$ gives $c \approx 0.81 < 1$.

6.3 Matrix Iteration

This polynomial machinery lifts directly to matrices. Let $M = I - A$. We maintain an approximate inverse $Y_n \approx M^{-1}$ and its residual $R_n = I - MY_n$. Starting from $Y_0 = I$ and $R_0 = A$, each step applies f to the residual:

$$Y_{n+1} = Y_n\cdot f(R_n), \qquad R_{n+1} = I - M\,Y_{n+1}. \tag{9}$$

Lemma 6.9 (General radix-kernel iteration)**.** *Under iteration (9):*

1. $R_n = E^{[n]}(A)$;
2. $R_n = O(A^{m^n})$;
3. $Y_n = \sum_{j=0}^{m^n-1} A^j + O(A^{m^n})$.

Proof. Base case. $R_0 = A = E^{[0]}(A)$. *Induction.* If $R_n = E^{[n]}(A)$, then $R_{n+1} = I - M\,Y_n f(R_n) = I - (I-R_n)f(R_n) = E(R_n) = E^{[n+1]}(A)$. Parts (2)–(3) follow from $E(z) = O(z^m)$ and $M\,Y_n = I - R_n = I - O(A^{m^n})$. □

6.4 Cost Analysis

This gain can be quantified directly and compared against prior methods. In the exact framework, one radix step consists of the kernel, concatenation, and a telescoping/power-update product. In the approximate framework (9), the

Table 1. Matrix product counts for $S_k(A)$ with $k \geq 10^p$. Cost per step is $\mu_m + 2$, minus 2 for the trivial first update ($Y_0 = I$) and the unnecessary last residual. Radix-24$^+$ uses radix-24 for all steps except possibly the last, where a smaller radix may be cheaper.

Method	Coeff.	10^3	10^6	10^9	10^{12}	10^{15}
Binary [6]	2.00	18	38	58	78	98
Quinary [5]	1.72	18	34	50	70	86
Radix-9	1.58	18	33	48	63	78
Radix-15	1.54	16	34	46	64	76
Radix-24	1.53	19	33	47	61	75
Radix-24$^+$	—	**14**	**29**	**44**	**59**	**75**

last operation is instead a residual recomputation. Each step still requires μ_m products for $f(R_n)$, 1 for $Y_{n+1} = Y_n \cdot f(R_n)$, and 1 for $R_{n+1} = I - M\,Y_{n+1}$: total $\mu_m + 2$ per step. Achieving accuracy $k = m^t$ requires $t = \log_m k$ steps:

$$(\mu_m + 2)\log_m k = \frac{\mu_m + 2}{\log_2 m}\log_2 k. \tag{10}$$

Corollary 6.10 (Radix-15 and radix-24 coefficients)**.** *An approximate radix-15 kernel with* $\mu_{15} = 4$ *achieves coefficient* $6/\log_2 15 \approx 1.54$. *An approximate radix-24 kernel with* $\mu_{24} = 5$ *achieves coefficient* $7/\log_2 24 \approx 1.53$.

The certified radix-15 and radix-24 kernels satisfy the prefix equations exactly at the parameter level, so the coefficient model applies to the mathematical iteration itself. Section 8 then checks that floating-point implementations exhibit the predicted m^n prefix growth and the expected product-count savings.

Table 1 compares all methods against prior work. Under naive kernel computation, higher radices were counterproductive and left coefficients at or above $2\log_2 k$. Once low-product kernels are available, the ordering reverses and higher radices become the asymptotically best options.

Since any step can use any available kernel, the optimal strategy uses radix-24 for all but the final step, where a smaller radix may suffice: e.g. for $k \geq 10^3$, the sequence $(24, 24, 2)$ gives $k = 1152$ in just 14 products. At $k = 10^{15}$, this saves 23 products over binary (75 vs 98, a 23% reduction).

The optimization principle: choose m to minimize $C(m)/\log_2 m$. For our kernels: $4/\log_2 5 \approx 1.72$ (quinary), $5/\log_2 9 \approx 1.58$ (radix-9), $6/\log_2 15 \approx 1.54$ (radix-15), $7/\log_2 24 \approx 1.53$ (radix-24). Higher radices help only if kernel cost grows sublinearly in $\log_2 m$.

6.5 Special Cases

The framework recovers known methods as special cases.

Hensel lifting ($m = 2$). $f(z) = 1 + z$, $E(z) = z^2$, $\mu_2 = 0$: $f(R_n) = I + R_n$ requires no GEMM once R_n is already available, so $Y_{n+1} = Y_n(I + R_n) = Y_n(2I - MY_n)$, coefficient $2/\log_2 2 = 2.00$.

Exact radix-m. $f = T_m$, $E(z) = z^m$, $R_n = A^{m^n}$: recovers the standard telescoping update.

The framework resolves both obstacles by applying the kernel to the residual, not to powers of A. The cost per step is $\mu_m + 2$—identical to the exact case—and the prefix grows as m^n regardless of spillover. Hensel lifting and exact telescoping emerge as special cases, so the framework subsumes all prior methods. With radix-24 it achieves coefficient 1.53, the smallest asymptotic coefficient among the kernels constructed here. The cost analysis assumes convergence; the next section examines when spillover preserves it.

7. Stability Analysis

Section 6 showed that the residual satisfies $R_n = E^{[n]}(A)$, where

$$E(z) = 1 - (1 - z)f(z)$$

is the kernel error map. Prefix growth alone does not imply bounded iterates: once spillover is present, the relevant question is whether repeated composition of E drives residuals to zero on the matrix set of interest. This section makes that question quantitative.

7.1 Error-Map Coefficients

Let

$$f(z) = \sum_{j=0}^{d} a_j z^j, \qquad a_0 = \cdots = a_{m-1} = 1,$$

so that the prefix is exact through degree $m - 1$. Then

$$E(z) = \sum_{j=m}^{d+1} e_j z^j, \tag{11}$$

with

$$e_m = 1 - a_m, \qquad e_j = a_{j-1} - a_j \ (m < j \leq d), \qquad e_{d+1} = a_d.$$

The leading coefficient e_m is only the first spillover datum; all of (11) matters for stability.

7.2 Scalar Tests

Define

$$q_\rho := \frac{1}{\rho}\max_{|z|=\rho}|E(z)| = \frac{1}{\rho}\max_{|z|\leq\rho}|E(z)|, \qquad \eta_\rho := \sum_{j=m}^{d+1}|e_j|\rho^{j-1},$$

and

$$s_0 := 0, \qquad s_n := 1 + m + \cdots + m^{n-1} = \frac{m^n - 1}{m - 1} \quad (n \geq 1).$$

Basin test. For a scalar state $\lambda_{n+1} = E(\lambda_n)$, convergence requires $E^{[n]}(\lambda) \to 0$. This is the exact scalar criterion.

Disk test. If $q_\rho \le 1$, then the disk $|z| \le \rho$ is forward invariant. Since $E(z) = z^m h(z)$ has a zero of order m at the origin, the generalized Schwarz lemma gives

$$|E^{[n]}(z)| \le \rho\, q_\rho^{s_n} \Big(\frac{|z|}{\rho}\Big)^{m^n}, \qquad |z| \le \rho. \tag{12}$$

Coefficient surrogate. The cheaper quantity η_ρ satisfies $q_\rho \le \eta_\rho$. Hence $\eta_\rho \le 1$ is a sufficient, though possibly conservative, way to verify the disk test.

7.3 A Functional-Calculus Theorem

The scalar bound (12) lifts immediately to any matrix class that admits a polynomial functional-calculus estimate.

Theorem 7.1 (Abstract matrix bound). *Let $\Omega \subset \mathbb{C}$ be compact, and suppose*

$$\sigma(A) \subset \Omega \subset \{z : |z| \le \rho\}, \qquad \|p(A)\|_2 \le M_\Omega \max_{z\in\Omega} |p(z)|$$

for every polynomial p. Set $r_\Omega := \max_{z\in\Omega} |z|$. If $q_\rho \le 1$, then for every $n \ge 0$,

$$\|R_n\|_2 \le M_\Omega\, \rho\, q_\rho^{s_n} \Big(\frac{r_\Omega}{\rho}\Big)^{m^n}.$$

Consequently,

$$\frac{\|(I-A)^{-1} - Y_n\|_2}{\|(I-A)^{-1}\|_2} \le M_\Omega\, \rho\, q_\rho^{s_n} \Big(\frac{r_\Omega}{\rho}\Big)^{m^n}.$$

Proof. Lemma 6.9 gives $R_n = E^{[n]}(A)$. Applying the assumed functional-calculus estimate to the polynomial $E^{[n]}$ and then using (12) on Ω yields the residual bound. Since $(I-A)^{-1} - Y_n = (I-A)^{-1} R_n$, the relative inverse-error estimate follows immediately. □

Corollary 7.2 (Normal matrices). *If A is normal, then Theorem 7.1 applies with $\Omega = \sigma(A)$ and $M_\Omega = 1$. In particular,*

$$\|R_n\|_2 \le \rho\, q_\rho^{s_n} \Big(\frac{\rho(A)}{\rho}\Big)^{m^n},$$

whenever $\rho \ge \rho(A)$ and $q_\rho \le 1$.

Corollary 7.3 (Numerical-range bound). *Let*

$$W(A) = \{x^* A x : \|x\|_2 = 1\}, \qquad w(A) = \max_{z\in W(A)} |z|.$$

By Crouzeix and Palencia [13],

$$\|p(A)\|_2 \le C_{\mathrm{Cr}} \max_{z\in W(A)} |p(z)|, \qquad C_{\mathrm{Cr}} \le 1 + \sqrt{2}.$$

Hence, if $\rho \ge w(A)$ and $q_\rho \le 1$, then

$$\|R_n\|_2 \le C_{\mathrm{Cr}}\, \rho\, q_\rho^{s_n} \Big(\frac{w(A)}{\rho}\Big)^{m^n}.$$

If moreover $\rho < 1$, then

$$\|(I-A)^{-1} - Y_n\|_2 \le \frac{C_{\mathrm{Cr}}}{1-\rho}\, \rho\, q_\rho^{s_n} \Big(\frac{w(A)}{\rho}\Big)^{m^n}.$$

Limitations and optional refinement. For the present certified kernels, the admissible radii remain below 1:

$$\rho_{q,15}^{\mathrm{cert}} \approx 0.971, \qquad \rho_{q,24}^{\mathrm{cert}} \approx 0.984.$$

The Crouzeix–Palencia constant is universal and therefore conservative, so tighter matrix-specific functional-calculus bounds could enlarge the usable region. The pseudospectral estimate below records one optional route in that direction, but it is not used in Section 8.

Proposition 7.4 (Pseudospectral refinement). *Let Γ_ε be a rectifiable contour enclosing $\sigma(A)$ such that*

$$\|(zI - A)^{-1}\|_2 \le \varepsilon^{-1} \qquad \text{for all } z \in \Gamma_\varepsilon,$$

and let $L_\varepsilon = \mathrm{len}(\Gamma_\varepsilon)$. Then for every polynomial p,

$$\|p(A)\|_2 \le \frac{L_\varepsilon}{2\pi\varepsilon} \max_{z\in\Gamma_\varepsilon} |p(z)|.$$

If $\Gamma_\varepsilon \subset \{z : |z| \le \rho\}$ and $r_\varepsilon := \max_{z\in\Gamma_\varepsilon} |z|$, then

$$\|R_n\|_2 \le \frac{L_\varepsilon}{2\pi\varepsilon}\, \rho\, q_\rho^{s_n} \Big(\frac{r_\varepsilon}{\rho}\Big)^{m^n}.$$

Proof. The contour bound is the standard polynomial functional calculus on Γ_ε; see, for example, [3, Chapter 17] or [14]. The final estimate is another application of Theorem 7.1. □

7.4 Certified Radii for Radix 15

For the certified radix-15 kernel, the error map summarized in Appendix 4 is centered at

$$\begin{aligned} E_{15}(z) = {} & 0.81479100\, z^{15} \\ & - 0.27315103\, z^{16} \\ & + 0.45836003\, z^{17}. \end{aligned}$$

Here the coefficient surrogate is essentially sharp:

$$\begin{aligned} \eta_{15}(\rho) = {} & 0.81479100\, \rho^{14} \\ & + 0.27315103\, \rho^{15} \\ & + 0.45836003\, \rho^{16}. \end{aligned}$$

The certified coefficient radius is

$$\rho_{\eta,15} = 0.970897959\ldots,$$

and the certified direct circle-test radius is

$$\rho_{q,15}^{\mathrm{cert}} = 0.970699488\ldots.$$

At the midpoint coefficients, midpoint circle maximization gives the estimated value

$$\rho_{q,15}^{\mathrm{est}} = 0.970897959\ldots,$$

which coincides with $\rho_{\eta,15}$ to the displayed digits. Thus, for practical purposes, the coefficient surrogate is sharp for radix 15.

At $w(A) = \rho = 0.95$, Corollary 7.3 yields

$$\|R_n\|_2 \le C_{\mathrm{Cr}}\, 0.95\, (0.725638)^{(15^n-1)/14},$$

so

$$\|R_2\|_2 \lesssim 1.36 \times 10^{-2}, \qquad \|R_3\|_2 \lesssim 6.3 \times 10^{-34}.$$

7.5 Certified Radii for Radix 24

For the certified radix-24 kernel, the error map is centered at

$$\begin{aligned} E_{24}(z) = {} & 0.35517021\, z^{24} + 0.43096456\, z^{25} + 0.16540498\, z^{26} \\ & - 0.13946877\, z^{27} - 0.24168001\, z^{28} - 0.10159049\, z^{29} \\ & + 0.10987210\, z^{30} + 0.20681982\, z^{31} + 0.15756340\, z^{32} \\ & + 0.05694420\, z^{33}. \end{aligned}$$

The coefficient surrogate certifies

$$\rho_{\eta,24} = 0.974559738\ldots,$$

while the interval-based direct circle test certifies the larger radius

$$\rho_{q,24}^{\mathrm{cert}} = 0.984176086\ldots.$$

At the midpoint coefficients, midpoint circle maximization gives the estimated value

$$\rho_{q,24}^{\mathrm{est}} = 0.984432110\ldots.$$

The gap between $\rho_{\eta,24}$ and $\rho_{q,24}^{\mathrm{cert}}$ is already visible in the certified intervals, so cancellations matter substantially for radix 24.

7.6 Summary

The error-map viewpoint separates three levels of statement:

- **exact scalar criterion:** $E^{[n]}(\lambda) \to 0$ for each scalar state of interest;
- **certified scalar tests:** the direct circle test $q_\rho \le 1$ or the cheaper coefficient surrogate $\eta_\rho \le 1$;
- **matrix-level corollaries:** normal, numerical-range, and pseudospectral bounds through Theorem 7.1.

For radix 15, the coefficient surrogate is effectively sharp and certifies numerical-range radii near 0.971. For radix 24, the direct circle test is materially better, certifying radii near 0.984 while the coefficient surrogate stops near 0.975. Section 8 compares these theorem predictions with observed matrix residuals.

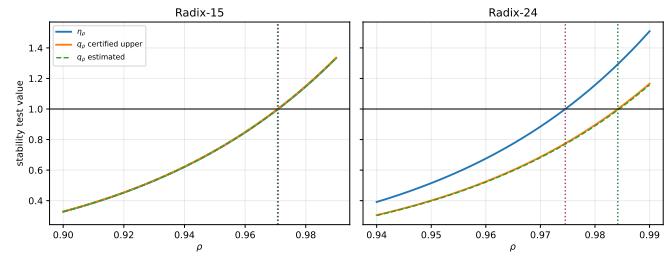

Figure 1. Certified scalar stability curves for the radix-15 and radix-24 kernels. The plotted data and certified crossing radii are included in the supplementary materials.

8. Numerical Experiments

The experiments play two roles. First, they validate the theorem-driven stability picture from Section 7: the certified radii, the gap between η_ρ and q_ρ, and the effect of nonnormality on matrix bounds. Second, they confirm that the certified higher-radix kernels still deliver the predicted product-count savings in illustrative inverse-approximation tests under the dense-matrix GEMM-count model.

Reproducibility. All tables and figures in this section are derived from the supplementary certification and validation data. All computations use Python with NumPy/SciPy and IEEE 754 double precision.

8.1 Certified Radius Curves

Figure 1 plots the coefficient surrogate η_ρ, the certified direct circle test for q_ρ, and midpoint circle maximization for the certified radix-15 and radix-24 kernels. The picture matches the theory from Section 7: for radix 15 the coefficient surrogate is essentially sharp, whereas for radix 24 the direct circle test is noticeably better.

The certified radii are

$$\rho_{\eta,15} = 0.970897959, \qquad \rho_{q,15}^{\mathrm{cert}} = 0.970699488,$$

and

$$\rho_{\eta,24} = 0.974559738, \qquad \rho_{q,24}^{\mathrm{cert}} = 0.984176086.$$

For radix 24, midpoint circle maximization raises the estimated crossing to 0.984432110, so the certified direct circle test captures most of the available cancellation already.

8.2 Normal and Nonnormal Validation

Figure 2 compares actual residual norms with the numerical-range bound from Corollary 7.3 on two small families. The first family is normal, with diagonal matrices at radii just inside the certified threshold. The second family is nonnormal, using 2×2 upper-triangular matrices with

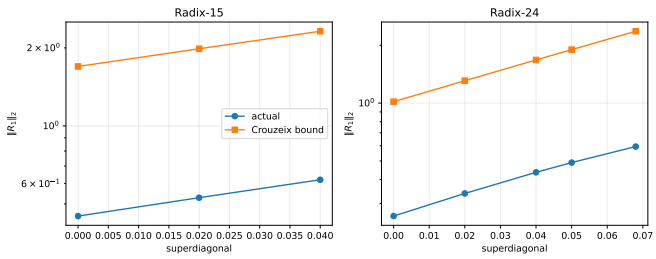

Figure 2. Observed residual norms versus the numerical-range bound on normal and nonnormal test families. The bound remains valid and becomes looser as nonnormality increases, as expected from the Crouzeix factor.

Table 2. Product counts for $S_k(A)$ at pure powers (binary baseline).

k	Binary	Quinary	Radix-9	Radix-15	Radix-24	Savings
$125 = 5^3$	14	**12**	—	—	—	14%
$625 = 5^4$	20	**16**	—	—	—	20%
$81 = 9^2$	14	—	**10**	—	—	29%
$729 = 9^3$	20	—	**15**	—	—	25%
$225 = 15^2$	16	—	—	**12**	—	25%
$3375 = 15^3$	24	—	—	**18**	—	25%
$576 = 24^2$	20	—	—	—	**14**	30%
$13824 = 24^3$	28	—	—	—	**21**	25%

fixed diagonal and increasing superdiagonal. These are illustrative validation cases for the bounds, not application-scale performance tests. The corresponding validation data are included in the supplementary materials.

Two qualitative points stand out. First, the certified bounds are conservative but correctly ordered: the actual residual decays far below the bound in the normal family, then moves closer to it as the matrix becomes more nonnormal. Second, the radius tests are not cosmetic. Supplementary scalar tests beyond the certified threshold show that crossing the certified boundary materially worsens the one-step residual even when later compositions may still decay for a particular scalar point.

8.3 *Product Counts and Approximate Inversion*

Stability certification matters only if the faster kernels still reduce work in the stated cost model, so we close by returning to the original GEMM-count question. Table 2 lists the exact matrix-product counts for pure powers. The higher-radix certified kernels still deliver the same asymptotic improvements: radix 9 reduces the binary constant from 2.00 to 1.58, radix 15 reaches 1.54, and radix 24 reaches 1.53.

Figure 3 keeps the original convergence comparison on a normal test family, while Figure 4 records the asymptotic coefficients themselves. These figures provide a secondary layer of evidence: once the certified kernels are available and stable on the relevant matrix set, the higher radix translates directly into fewer matrix products under the adopted cost model.

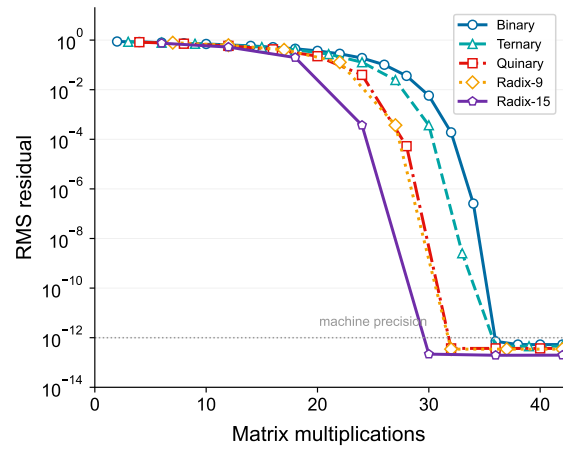

Figure 3. Convergence comparison on a normal family. Higher-radix methods reach machine precision with fewer matrix products once the certified stability conditions hold.

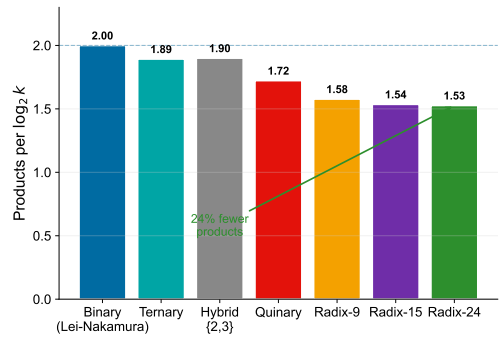

Figure 4. Asymptotic coefficients (matrix products per $\log_2 k$). Among the kernels studied here, the certified radix-24 kernel attains the smallest asymptotic coefficient, 1.53.

Summary. The experiments line up with the theory in three layers. The scalar curves explain which radii are certified. The normal and nonnormal illustrative matrix tests validate the matrix bounds. The remaining convergence and product-count plots show that once those stability conditions are met, the certified higher-radix kernels realize the expected 25–30% reduction in matrix products.

9. Conclusion

We set out to extend Neumann series evaluation beyond the known radix-5 kernel [5], where no higher-radix constructions existed and where nonexact higher-radix kernels lacked a matrix-level convergence theory. We addressed this gap with four linked contributions:

1. The radix-9 kernel achieves $\mu_9 = 3$ with rational coefficients. It is the first exact construction beyond radix 5 and yields $5 \log_9 k \approx 1.58 \log_2 k$ products, 21% fewer than binary splitting.
2. Certified approximate radix-15 and radix-24 kernels, obtained by numerical discovery plus interval certification, yield $6/\log_2 15 \approx 1.54$ and $7/\log_2 24 \approx 1.53$.
3. A general radix-kernel framework (Section 6) applies kernels to residuals, enabling iterative use of approximate kernels despite spillover.
4. Matrix-level convergence follows from an abstract theorem with normal, numerical-range, and pseudospectral corollaries that quantify residual decay for nonnormal matrices.

The transition from exact (radix-9) to certified approximate (radix-15, radix-24) kernels suggests that spillover is the norm at higher radices. Our local structure analysis supports that interpretation without claiming a global impossibility theorem: near the certified solutions, the prefix-exact set is a positive-dimensional manifold and the leading-spillover cancellation imposes one additional local constraint.

Supplementary materials record the certified radii, interval enclosures, and validation data underlying the numerical sections. Those materials make the paper's claims auditable; companion resources are available through the project GitHub repository and the tutorial site.

Several mathematical directions remain:

1. *Exact vs. approximate boundary*: Can we prove that exact radix-15 kernels do not exist, or find one?
2. *Higher radices*: Can approximate kernels with μ_m growing sublogarithmically improve the coefficient beyond 1.53?
3. *Structured matrices*: How do these algorithms extend to sparse or Toeplitz settings?

Practical runtime measurement, larger-scale matrix tests, and mixed-precision or parallel/GPU implementations remain future work. Those questions matter, but they are distinct from the present paper's goal of establishing low-product kernels, certified stability radii, and the resulting GEMM-count consequences.

More broadly, the paper reframes truncated Neumann-series evaluation as a kernel-design problem: better kernels translate directly into better series algorithms. That viewpoint may also benefit log-determinant estimation [4], where polynomial traces from radix-kernel sums could replace individual power traces while reducing the number of matrix products.

Acknowledgements

This work was supported by the Laboratory Directed Research and Development Program of Oak Ridge National Laboratory, managed by UT-Battelle, LLC, for the U.S. Department of Energy.

References

[1] Paul F. Dubois, Anne Greenbaum, and Garry H. Rodrigue. "Approximating the Inverse of a Matrix for Use in Iterative Algorithms on Vector Processors". In: *Computing* 22.3 (1979), pp. 257–268. DOI: 10.1007/BF02243566.

[2] Yousef Saad. *Iterative Methods for Sparse Linear Systems*. 2nd. SIAM, 2003. DOI: 10.1137/1.9780898718003.

[3] Nicholas J. Higham. *Functions of Matrices: Theory and Computation*. SIAM, 2008. DOI: 10.1137/1.9780898717778.

[4] Piyush Sao. "What Trace Powers Reveal About Log-Determinants: Closed-Form Estimators, Certificates, and Failure Modes". In: *arXiv preprint arXiv:2601.12612* (2025). URL: https://arxiv.org/abs/2601.12612.

[5] Oscar Gustafsson et al. "Approximate Neumann Series or Exact Matrix Inversion for Massive MIMO?" In: *Proc. IEEE Symposium on Computer Arithmetic (ARITH)*. 2017, pp. 62–63. DOI: 10.1109/ARITH.2017.14.

[6] Shao-Feng Lei and Tadao Nakamura. "On the Fast Evaluation of Parallel ARMA Filters and Related Matrix Power Sums". In: *IEEE Transactions on Circuits and Systems I* 39.4 (1992), pp. 299–300. DOI: 10.1109/81.129460.

[7] Vassil S. Dimitrov and Todor V. Cooklev. "Two Algorithms for Modular Exponentiation Using Nonstandard Arithmetics". In: *IEICE Transactions on Fundamentals* E78-A.1 (1995), pp. 82–87.

[8] David Westreich. "A Practical Method for Computing the Exponential of a Matrix and its Integral". In: *IEEE Transactions on Circuits and Systems* 36.1 (1989), pp. 162–164. DOI: 10.1109/31.16585.

[9] Yusuke Matsumoto, Naofumi Takagi, and Kazuyoshi Takagi. "Efficient Algorithms for Computing Power Sums Using Higher Radices". In: *IEICE Transactions on Fundamentals* E101-A.2 (2018), pp. 467–471. DOI: 10.1587/transfun.E101.A.467.

[10] Michael S. Paterson and Larry J. Stockmeyer. "On the Number of Nonscalar Multiplications Necessary to Evaluate Polynomials". In: *SIAM Journal on Computing* 2.1 (1973), pp. 60–66. DOI: 10.1137/0202007.

[11] Massimiliano Fasi. "Optimality of the Paterson–Stockmeyer Method for Evaluating Matrix Polynomials and Rational Matrix Functions". In: *Linear Algebra and its Applications* 574 (2019), pp. 182–200. DOI: 10.1016/j.laa.2019.04.001.

[12] Jorge Sastre and Javier Ibáñez. "Boosting the Computation of the Matrix Exponential". In: *Mathematics* 9.14 (2021), p. 1600. DOI: 10.3390/math9141600.

[13] Michel Crouzeix and Carlos Palencia. "The Numerical Range Is a $(1+\sqrt{2})$-Spectral Set". In: *SIAM Journal on Matrix Analysis and Applications* 38.2 (2017), pp. 649–655. DOI: 10.1137/17M1116672.

[14] Lloyd N. Trefethen and Mark Embree. *Spectra and Pseudospectra: The Behavior of Nonnormal Matrices and Operators*. Princeton University Press, 2005.

Appendix 1. Kernel Coefficients

This appendix provides display formulas for the kernels and supplementary numerical validation on test matrices. For the certified approximate kernels, the corresponding interval data are summarized in Appendix 4.

Appendix 1.1 Quinary Kernel (2 Products)

The quinary kernel $T_5(B) = I + B + B^2 + B^3 + B^4$ is computed using 2 matrix products:

$$U = B^2, \tag{13}$$

$$V = U(B + U) = B^3 + B^4, \tag{14}$$

$$T_5(B) = I + B + U + V. \tag{15}$$

This is exact: all 5 coefficients equal 1.

Appendix 1.2 Radix-9 Kernel (3 Products)

The radix-9 kernel $T_9(B) = \sum_{j=0}^{8} B^j$ uses 3 products:

$$U = B^2, \tag{16}$$

$$V = U(B + 2U) = B^3 + 2B^4, \tag{17}$$

$$P = (0.15B + 2U + V)(0.275B - 0.125U + 0.25V), \tag{18}$$

$$T_9(B) = I + B + \tfrac{767}{800}U + \tfrac{15}{32}V + P. \tag{19}$$

The rational coefficients in (19) are $\frac{767}{800} = 0.95875$ and $\frac{15}{32} = 0.46875$.

Appendix 1.3 Radix-9 Derivation Note

Write the third product in the power basis as

$$P = \sum_{j=1}^{4} p_j B^j, \qquad Q = \sum_{j=1}^{4} q_j B^j,$$

and then add low-degree corrections $xU + yV$. Matching $W = P \cdot Q$ to the target coefficients of $T_9(B)$ yields the linear coefficient conditions

$$\begin{aligned} p_1q_1 + x &= 1, \\ p_1q_2 + p_2q_1 + y &= 1, \\ p_1q_3 + p_2q_2 + p_3q_1 + 2y &= 1, \\ p_1q_4 + p_2q_3 + p_3q_2 + p_4q_1 &= 1, \\ p_2q_4 + p_3q_3 + p_4q_2 &= 1, \\ p_3q_4 + p_4q_3 &= 1, \\ p_4q_4 &= 1. \end{aligned}$$

Imposing the structural relations coming from $P = \alpha B + \beta U + V$ and $Q = \gamma B + \delta U + \varepsilon V$, namely $p_3 = 1$, $p_4 = 2$, and $q_4 = 2q_3$, closes the resulting linear system and yields the displayed rational solution

$$(\alpha, \beta, \gamma, \delta, \varepsilon, x, y) = \left(\tfrac{3}{20},\ 2,\ \tfrac{11}{40},\ -\tfrac{1}{8},\ \tfrac{1}{4},\ \tfrac{767}{800},\ \tfrac{15}{32}\right).$$

Appendix 1.4 Numerical Validation

We validate convergence on an illustrative 64×64 symmetric positive definite (SPD) matrix A with geometrically spaced eigenvalues in $[10^{-2}, 1]$, giving $\kappa(I - A) \approx 10^4$. In this appendix, let $B = I - A$ (distinct from the kernel argument B in the main text).

Setup. For series index k, each method computes an approximation S_k to $(I - B)^{-1}$. We measure the relative residual $\|I - (I - B)S_k\|_F / \sqrt{d}$, where $\|\cdot\|_F$ denotes the Frobenius norm.

Results. Table 3 lists the matrix product count required to reach machine precision ($\approx 10^{-13}$).

Table 3. Matrix product counts to reach 10^{-13} residual ($d = 64$, $\kappa \approx 10^4$).

Method	Products	Residual
Binary ($m = 2$)	38	5.6×10^{-13}
Ternary ($m = 3$)	36	4.7×10^{-13}
Quinary ($m = 5$)	32	3.9×10^{-13}
Radix-9	32	4.1×10^{-13}
Radix-15 (approx.)	30	4.2×10^{-13}
Radix-24 (approx.)	28	3.4×10^{-14}

In this illustrative dense-matrix test, radix-24 uses the fewest products and, with refined coefficients, reaches machine precision despite nonzero spillover. The naive method requires $k - 1 = 10^4$ products for comparable accuracy.

Efficiency comparison. Per radix step, the update cost is $C(m) = \mu_m + 2$ products:

- Binary: $C(2) = 0 + 2 = 2$, coefficient $\frac{2}{\log_2 2} = 2.00$;
- Ternary: $C(3) = 1 + 2 = 3$, coefficient $\frac{3}{\log_2 3} \approx 1.89$;
- Quinary: $C(5) = 2 + 2 = 4$, coefficient $\frac{4}{\log_2 5} \approx 1.72$;
- Radix-9: $C(9) = 3 + 2 = 5$, coefficient $\frac{5}{\log_2 9} \approx 1.58$;
- Radix-15: $C(15) = 4 + 2 = 6$, coefficient $\frac{6}{\log_2 15} \approx 1.54$;
- Radix-24: $C(24) = 5 + 2 = 7$, coefficient $\frac{7}{\log_2 24} \approx 1.53$.

Radix-9 achieves the best rate among exact kernels; radix-24 achieves the best overall.

Appendix 1.5 Approximate Kernels: Residual vs Radix Trade-off

For a single truncated Neumann computation (no recursion), we compare approximate kernels at fixed matrix product budgets (6 and 7 products). Test matrix: $d = 64$, geometrically spaced eigenvalues in $[0.01, 1]$, giving $\rho(B) = 0.99$.

Key finding: Higher radix dominates. Table 4 compares kernels at 6 products and 7 products budgets. Even with high kernel residual, higher radix gives better inverse approximation. Error is measured in spectral norm: $\|S - A^{-1}\|_2 / \|A^{-1}\|_2$, where $\|\cdot\|_2$ denotes the largest singular value.

Table 4. Inverse approximation error (spectral norm) for single kernel application. "vs Exact": ratio to exact Neumann at same radix.

Method	Prods	Kern. Res.	Inv. Err.	vs Exact
6-product budget				
Standard S_7 (radix 7)	–	0	0.932	–
Radix-15	4	5×10^{-16}	0.860	1.000
Radix-16	4	1×10^{-2}	0.852	1.000
Radix-17	4	1×10^{-1}	0.844	1.001
7-product budget				
Standard S_8 (radix 8)	–	0	0.923	–
Radix-24	5	2×10^{-15}	0.786	1.000
Radix-27	5	1×10^{-3}	0.762	1.000
Radix-31	5	1×10^{-1}	0.733	1.002
Radix-33	5	6×10^{-1}	0.722	1.007

Interpretation. The "vs Exact" column shows degradation from kernel approximation error. Even radix-33 with residual 0.6 is only 0.7% worse than an exact radix-33 kernel. The dominant factor is radix m: error scales as $\rho(B)^m$, so larger m gives exponentially better approximation, overwhelming the kernel residual effect.

Coefficients for radix-33 kernel (5 products). The approximate coefficients c_j for $\widetilde{T}_{33}(B) = \sum_{j=0}^{32} c_j B^j$: Coefficients

j	c_j values
0–10	1.00, 1.00, 1.00, 0.98, 1.00, 1.02, 1.00, 0.69, 1.12, 0.99, 1.05
11–21	0.88, 1.11, 0.69, 0.73, 1.01, 1.27, 1.04, 1.08, 1.06, 0.98, 1.00
22–32	0.99, 0.88, 1.13, 0.95, 1.02, 0.97, 1.01, 0.99, 1.00, 1.00, 1.00

range from 0.69 to 1.27, yet the inverse error is only 1% worse than exact. No spillover terms exist beyond degree 32.

Appendix 2. Radix-15 Kernel: Explicit Circuit

The circuit (6) with $P_1 = B^2$ is fully specified by the following coefficients, obtained by numerically refining an approximate solution of the bilinear ansatz. Each product $P_i = L_i \cdot R_i$; rows of the table give coefficients of $\{I, B, P_1, \ldots\}$. The table is rounded for display; the corresponding certified interval data are summarized in Appendix 4.

	I	B	P_1	P_2	P_3
L_2	0.237957446781864	0.241306784696238	1.573677404167023		
R_2	0.047776286323906	−0.046512791269008	0.888968295817947		
L_3	0.262503282946680	0.918866331886906	0.841863255429817	0.819355416606328	
R_3	0.183860759371110	0.068445979922426	0.133854440389316	−1.405242484443339	
L_4	−0.004559362003080	−1.384977065724292	0.021866758593494	0.122021888194117	0.274863153197116
R_4	−0.701388642425678	0.601773558132688	−0.624151559177423	−0.137295755918682	0.328424149078130

Final combination:

$$\widetilde{T}_{15} = I + 0.077578347530026\, B + 1.778433959744106\, P_1 + 0.663737059861237\, P_2 - 0.024081230623538\, P_3 + P_4.$$

At these coefficients, the prefix is exact to double precision: $\max_{0\le j\le 14} |[\widetilde{T}_{15}]_j - 1| < 5\times10^{-16}$. The spillover is $0.185\, B^{15} + 0.458\, B^{16} + O(B^{17})$, giving error map leading coefficient $c = 1 - [\widetilde{T}_{15}]_{15} \approx 0.81$.

Appendix 3. Radix-24 Kernel: Explicit Circuit

The 5-product circuit with $P_1 = B^2$ is specified by the coefficients below, obtained by the same type of numerical refinement. Each product $P_i = L_i \cdot R_i$. The coefficients are printed to 16 decimal places so that the displayed circuit itself reproduces the prefix equations to roundoff accuracy. The corresponding certified interval data are summarized in Appendix 4.

Vector	Basis element	Coefficient
L_2	I	0.3282879733298226
	B	0.3334609994248492
	P_1	-1.6413300182113680
R_2	I	0.1016599828103740
	B	0.6352737284711171
	P_1	0.9424901854962264
L_3	I	-0.4039110435614196
	B	-0.1853636671117334
	P_1	-0.0193103621489443
	P_2	0.7643808028090610
R_3	I	-0.6831517007752321
	B	0.0593700272785998
	P_1	0.9555951950613576
	P_2	-0.8352898821299267
L_4	I	0.8503961179081752
	B	-0.1700053748177661
	P_1	0.4007962270352739
	P_2	-0.7583911189268658
	P_3	-0.4202327816885083
R_4	I	0.3328223421061135
	B	0.8160093825695688
	P_1	-0.2722035796025236
	P_2	0.5557279222526658
	P_3	-0.4002282188877652
L_5	I	-0.0031207753630199
	B	0.2316286284512770
	P_1	-0.6242391015078156
	P_2	0.1290807812447923
	P_3	0.5768300240418512
	P_4	-0.5504031164707843
R_5	I	-0.5424304856178253
	B	-0.2443407079828936
	P_1	0.0247860011528895
	P_2	0.2849712559148055
	P_3	-0.0225503875904663
	P_4	-0.6711249224603717

Final combination:

$$\begin{aligned}\widetilde{T}_{24} = I &+ 0.9847461487306262\,B + 0.8611761979377541\,P_1 \\ &- 0.7865479053719371\,P_2 - 0.0121109383555383\,P_3 \\ &+ 0.4000630253949220\,P_4 + P_5.\end{aligned}$$

Evaluating the displayed 16-decimal circuit in high precision gives $\max_{0\le j\le 23} |[\widetilde{T}_{24}]_j - 1| < 3\times 10^{-16}$. The refined spillover begins at

$$\begin{aligned}&0.6448297923472228\,B^{24} + 0.2138652343604695\,B^{25} \\ &\quad + 0.0484602542811774\,B^{26} + 0.1879290230749536\,B^{27} \\ &\quad + 0.4296090302017557\,B^{28} + 0.5311995248359085\,B^{29} \\ &\quad + 0.4213274226205107\,B^{30} + 0.2145076021946625\,B^{31} \\ &\quad + 0.0569441977032723\,B^{32},\end{aligned}$$

giving error-map leading coefficient $c = 1 - [\widetilde{T}_{24}]_{24} \approx 0.3551702076527769$.

Appendix 4. Certified Approximate Kernels

The circuit tables in Appendix 2 and Appendix 3 give decimal representatives of the certified kernels. This appendix summarizes the supplementary interval data underlying the certified kernels.

Certification workflow. For each radix-m ansatz we form the prefix map

$$F_m(\theta) = ([f_\theta]_0 - 1, \ldots, [f_\theta]_{m-1} - 1).$$

At a numerically discovered approximate solution θ_0, we compute $DF_m(\theta_0)$ and select a well conditioned square subsystem. Holding the remaining variables fixed, we apply a Krawczyk inclusion test on a small nearby box in those variables. This certifies the existence of an exact root of the prefix equations near the approximate solution. Evaluating the polynomial and error map over that certified box then yields interval enclosures for the spillover and error-map coefficients. For radix 15, the active subsystem has 15 variables, condition estimate 3.83×10^1, and certified maximum box radius 4.5×10^{-14}. For radix 24, the corresponding values are 24 variables, 2.38×10^2, and 6.2×10^{-14}.

Display precision policy. Appendix circuit tables display decimal midpoint coefficients. The radix-24 table is printed to 16 decimal places; evaluating those printed decimals gives prefix defect below 3×10^{-16}. Certified stability claims always use interval data rather than rounded display coefficients. The complete interval records, active sets, and radii are provided in the supplementary certification files `certified_coeffs_radix15.json` and `certified_coeffs_radix24.json`.

Certified kernels. For radix 15, the verified box radius is below 4.5×10^{-14} and the certified error map satisfies

$$E_{15}(z) = e_{15}z^{15} + e_{16}z^{16} + e_{17}z^{17},$$

with intervals centered at

$$\begin{aligned}e_{15} &= 0.81479100,\\ e_{16} &= -0.27315103,\\ e_{17} &= 0.45836003.\end{aligned}$$

For radix 24, the verified box radius is below 6.2×10^{-14} and the certified error map coefficients are centered at

$$\begin{aligned}e_{24} &= 0.35517021, & e_{25} &= 0.43096456,\\ e_{26} &= 0.16540498, & e_{27} &= -0.13946877,\\ e_{28} &= -0.24168001, & e_{29} &= -0.10159049,\\ e_{30} &= 0.10987210, & e_{31} &= 0.20681982,\\ e_{32} &= 0.15756340, & e_{33} &= 0.05694420.\end{aligned}$$

The full certified intervals, including radii, are part of the supplementary certification data cited above.

Certified radii. The coefficient-surrogate radii are certified at

$$\rho_{\eta,15} = 0.970897959\ldots, \qquad \rho_{\eta,24} = 0.974559738\ldots.$$

The direct circle-maximization test with coefficient intervals certifies

$$\rho^{\mathrm{cert}}_{q,15} = 0.970699488\ldots, \qquad \rho^{\mathrm{cert}}_{q,24} = 0.984176086\ldots.$$

For comparison, midpoint circle maximization gives the estimated values $0.970897959\ldots$ and $0.984432110\ldots$.

Appendix 5. Local Structure of the Prefix-Exact Set

The certification step yields concrete base points at which the prefix maps have full row rank. This supports a local structural theorem for the approximate-kernel ansätze.

Proposition Appendix 5.1 (Local structure near the certified kernels)**.** *Let F_m denote the prefix map for one of the certified bilinear ansätze, with parameter count p. At the certified radix-15 and radix-24 base points, the Jacobian DF_m has full row rank m. Consequently, the prefix-exact set $F_m(\theta) = 0$ is locally a smooth manifold of dimension $p - m$: 13 for radix 15 and 17 for radix 24. If the leading spillover coefficient is appended as an additional constraint, the local dimension drops by one.*

Proof. This is again a computer-assisted rank verification. At the certified base points, the Jacobian matrices are verified to have full row rank. Full row rank of DF_m invokes the implicit function theorem. Appending one independent spillover constraint increases the rank by one for both certified kernels, so the local dimension decreases by one. □

Remark Appendix 5.2. This is a local statement only. It does not prove the global nonexistence of exact radix-15 or radix-24 kernels, nor does it exclude exact kernels in other ansätze.

Appendix 6. Companion Software and Reproducibility

Supplementary materials accompany the paper and include the certified interval data and validation data used in the numerical sections. High-level companion resources are available through the project GitHub repository and the tutorial site.