RESEARCH PREPRINT

Fast Evaluation of Truncated Neumann Series by Low-Product Radix Kernels

Piyush Sao

Oak Ridge National Laboratory, Oak Ridge, TN, USA

Abstract

Truncated Neumann series $S_k(A) = I + A + \cdots + A^{k-1}$ are used in approximate matrix inversion and polynomial preconditioning. In dense settings, matrix-matrix products dominate the cost of evaluating S_k. Naive evaluation needs $k-1$ products, while splitting methods reduce this to $O(\log k)$. Repeated squaring, for example, uses $2\log_2 k$ products, so further gains require higher-radix kernels that extend the series by m terms per update. Beyond the known radix-5 kernel, explicit higher-radix constructions were not available, and the existence of exact rational kernels was unclear.

We construct *radix kernels* for $T_m(B) = I + B + \cdots + B^{m-1}$ and use them to build faster series algorithms. For radix 9, we derive an *exact* 3-product kernel with rational coefficients, which is the first exact construction beyond radix 5. This kernel yields $5\log_9 k = 1.58\log_2 k$ products, a 21% reduction from repeated squaring. For radix 15 and radix 24, numerical discovery followed by interval certification produces *certified approximate* kernels whose prefixes are exact and whose spillover coefficients are enclosed by machine-readable interval files. Because spillover breaks the standard telescoping update, we introduce a residual-based radix-kernel framework that accommodates approximate kernels and retains coefficient $(\mu_m + 2)/\log_2 m$. We then prove a matrix-level convergence theorem that specializes to the normal case, a numerical-range bound via Crouzeix's theorem, and a pseudospectral refinement for strongly nonnormal matrices. The certified kernels yield update coefficients $6/\log_2 15 \approx 1.54$ and $7/\log_2 24 \approx 1.53$, while the associated stability tests certify numerical-range radii near 0.971 and 0.984. Theory-driven experiments validate the decay bounds, quantify the gap between cheap coefficient tests and direct circle maximization, and confirm the expected product-count savings.

Keywords: Neumann series; matrix polynomial; radix kernel; matrix multiplication; computational complexity

1. Introduction

The truncated Neumann series is a core object in numerical linear algebra. For a matrix A with spectral radius $\rho(A) < 1$, it is defined by

$$S_k(A) = \sum_{j=0}^{k-1} A^j = I + A + A^2 + \cdots + A^{k-1}. \tag{1}$$

As $k \to \infty$, $S_k(A)$ converges to $(I - A)^{-1}$, so finite truncations give practical approximate inverses. These approximations appear in polynomial preconditioning [1, 2], resolvent approximation [3], log-determinant estimation [4], and massive MIMO signal detection [5]. In dense problems, the dominant cost is matrix-matrix multiplication (GEMM), so reducing product count is the main efficiency goal.

Existing splitting methods already reduce the naive $k-1$ products to $O(\log k)$ [6]. Binary splitting (repeated squaring) uses $2\log_2 k$ products. Hybrid radix-$\{2, 3\}$ schemes [7] lower the constant only slightly, to about $1.9\log_2 k$. Thus asymptotic order is no longer the bottleneck; the leading coefficient is.

Improving that coefficient therefore requires better radix kernels. Let $T_m(B) = I + B + \cdots + B^{m-1}$ and let μ_m denote the minimum number of products to evaluate T_m. The decomposition

$$S_{mn}(A) = S_n(A) \cdot T_m(A^n). \tag{2}$$

shows that one radix-m update costs $C(m) = \mu_m + 2$ products: μ_m for the kernel, one for concatenation, and one for updating the next power. Hence the asymptotic coefficient is $C(m)/\log_2 m$. Each product can at most double polynomial degree, so $\mu_m \geq \lceil \log_2(m-1) \rceil$. This lower bound gives $\mu_9 \geq 3$ and $\mu_{15} \geq 4$, which suggests that higher-radix kernels could improve constants if such kernels exist. Gustafsson et al. [5] achieved $\mu_5 = 2$ for radix

This manuscript has been authored by UT-Battelle, LLC under Contract No. DE-AC05-00OR22725 with the U.S. Department of Energy. The publisher, by accepting the article for publication, acknowledges that the United States Government retains a non-exclusive, paid-up, irrevocable, world-wide license to publish or reproduce the published form of this manuscript, or allow others to do so, for United States Government purposes. The Department of Energy will provide public access to these results of federally sponsored research in accordance with the DOE Public Access Plan (http://energy.gov/downloads/doe-public-access-plan).

5, but explicit constructions beyond radix 5 were missing. This leaves two linked questions: can radix 9 reach its lower bound exactly, and can higher-radix kernels remain useful when they are only approximate?

We answer these questions with four contributions. First, we construct an exact radix-9 kernel with $\mu_9 = 3$ and rational coefficients (Theorem 4.1), giving $5\log_9 k = 1.58\log_2 k$ products. This is the first exact construction beyond radix 5 and gives a 21% improvement over binary splitting. Second, for radices 15 and 24 we combine numerical discovery with interval certification to obtain prefix-exact higher-radix kernels with controlled spillover. Third, we develop a residual-based general radix-kernel framework that accommodates spillover while preserving coefficient $(\mu_m + 2)/\log_2 m$. Fourth, we prove matrix-level stability bounds for that framework, including normal, numerical-range, and pseudospectral corollaries, and validate them numerically. Within this framework, radix 15 reaches $6/\log_2 15 \approx 1.54$ and radix 24 reaches $7/\log_2 24 \approx 1.53$, the best known asymptotic coefficient.

The paper proceeds as follows. Section 2 reviews background and the series-specific prior work. Section 3 places the kernel viewpoint in the broader polynomial-evaluation literature. Section 4 presents the radix-9 construction. Section 5 separates numerical search from certified approximate higher-radix kernels. Section 6 develops the general framework and its cost analysis. Section 7 develops matrix-level stability theorems, and Section 8 validates those theorems numerically while comparing product counts.

2. Background

We review existing methods and show that naive kernel computation is their main limitation. Given $A \in \mathbb{R}^{d\times d}$ with $\rho(A) < 1$, define the truncated Neumann series as in (1):

$$S_k(A) = \sum_{j=0}^{k-1} A^j = I + A + A^2 + \cdots + A^{k-1}. \tag{3}$$

We measure cost by counting matrix products (GEMMs), which dominate when $d \gtrsim 100$ for dense matrices.

2.1 Splitting Methods

Several radix-based schemes exist, differing in how many terms they add per step.

- **Naive evaluation.** The naive approach accumulates powers $I, A, A^2, \ldots$, requiring $k-1$ matrix products—prohibitive for large k.
- **Radix-m splitting.** Faster algorithms use the factorization $S_{mn}(A) = S_n(A)\cdot(I + A^n + \cdots + A^{(m-1)n})$. A radix-$m$ scheme iterates this identity. Each step computes:
  - powers $A^{2n}, \ldots, A^{(m-1)n}$ from A^n ($m-2$ products),
  - the product $S_n \cdot (I + A^n + \cdots + A^{(m-1)n})$ (1 product),
  - and $A^{mn} = A^{(m-1)n}\cdot A^n$ (1 product).

  The naive cost per step is $C(m) = m$ products, giving coefficient $m/\log_2 m$.
- **Binary splitting.** For $m = 2$, the identity simplifies to $S_{2n} = S_n(I + A^n)$, using only 2 products per step [8, 6]: (i) $A^{2n} = A^n \cdot A^n$; (ii) $S_{2n} = S_n \cdot (I + A^n)$. Total: $2\log_2 k$ products.
- **Ternary splitting.** For $m = 3$, $S_{3n} = S_n(I + A^n + A^{2n})$: (i) $A^{2n} = A^n \cdot A^n$; (ii) $A^{3n} = A^{2n}\cdot A^n$; (iii) $S_{3n} = S_n \cdot (I + A^n + A^{2n})$. Total: $3\log_3 k \approx 1.89\log_2 k$ products.

Optimal radix for naive splitting. The coefficient $C(m)/\log_2 m$ equals $m/\log_2 m$, minimized at $m = e \approx 2.718$; thus $m = 3$ is optimal among integers:

Radix m	Cost $C(m)$	Coefficient $C(m)/\log_2 m$
2	2	$2/\log_2 2 = 2.00$
3	3	$3/\log_2 3 \approx 1.89$
4	4	$4/\log_2 4 = 2.00$

Thus ternary has the best coefficient among pure radix methods. Dimitrov and Cooklev [7] extended this to hybrid radix-$\{2,3\}$ methods for arbitrary k, achieving coefficient ≈ 1.9. These mixed $\{2,3\}$ strategies are near-optimal, but only with naive kernel computation.

2.2 Radix Kernels

The splitting methods above compute the partial sum $T_m(B) = I + B + \cdots + B^{m-1}$ naively, requiring $m-2$ products for powers $B^2, \ldots, B^{m-1}$. We call T_m the *kernel*. If the kernel cost μ_m grew only logarithmically in m, the coefficient $(\mu_m + 2)/\log_2 m$ could fall below 1.89. Two identities make this possible:

- **Block concatenation.** $S_{mn}(A) = S_n(A)\cdot T_m(A^n)$ for any $m, n \geq 1$, verified by expanding both sums.
- **Telescoping.** $T_m(B)(I - B) = I - B^m$ (the geometric series identity), which extracts B^m from the kernel without extra products.

Cost model. Let μ_m denote the minimum number of products to compute $T_m(B)$. A radix-m update $(S_n, A^n) \to (S_{mn}, A^{mn})$ uses:

- μ_m products for the kernel,
- 1 product for telescoping: $I - A^{mn} = T_m(A^n)(I - A^n)$,
- 1 product for concatenation: $S_{mn} = S_n \cdot T_m(A^n)$.

The total is $C(m) = \mu_m + 2$, giving coefficient $C(m)/\log_2 m$.

Lower bound. Since each product at most doubles polynomial degree, p products starting from $\{I, B\}$ reach degree at most 2^p. Because T_m has degree $m-1$, evaluating it requires at least $\lceil \log_2(m-1) \rceil$ products. In particular, $\mu_5 \geq 2$ and $\mu_9 \geq 3$. Can these bounds be achieved?

2.3 Prior Work: Quinary (Radix-5) Kernel

For radix-5, the answer is yes. Gustafsson et al. [5] achieved $\mu_5 = 2$ with the following construction: (i) $U := B^2$; (ii) $V := U(B+U) = B^3 + B^4$; (iii) $T_5(B) := I + B + U + V$. Direct verification confirms $T_5(B) = I + B + B^2 + B^3 + B^4$.

With $\mu_5 = 2$, a quinary update costs $C(5) = 4$ products, giving coefficient $4/\log_2 5 \approx 1.72$.

Open questions. The quinary kernel reduces the coefficient, but gaps remain. First, no radix-9 construction exists, though the degree-doubling bound permits $\mu_9 = 3$, giving coefficient $5/\log_2 9 \approx 1.58$. Second, it is unknown whether exact rational kernels exist for higher radices such as $m = 15$. Approximate constructions may have nonzero *spillover* into degrees $\geq m$, breaking the telescoping identity (2.2) and preventing exact extraction of B^m. No framework exists for such kernels in series algorithms. These gaps set the agenda for the rest of the paper. Section 3 places them in the broader matrix-polynomial-evaluation literature. Sections 4–6 then develop exact and approximate higher-radix kernels, and the later sections analyze their stability and numerical performance.

3. Related Work

The background above covered splitting methods specific to truncated Neumann series. A second line of context comes from general matrix-polynomial evaluation, where special structure can justify product counts that generic schemes cannot reach.

Within Neumann-series evaluation itself, Westreich [8] obtained $3\log_2 k$ products by binary splitting, and Lei and Nakamura [6] reduced this to $2\log_2 k - 1$ by reusing intermediate computations. Dimitrov and Cooklev [7] then introduced hybrid radix-$\{2,3\}$ methods and, under naive kernel evaluation, argued that higher radices offered no practical improvement. Matsumoto, Takagi, and Takagi [9] overturned that pessimistic asymptotic conclusion by showing that higher radices can improve the leading coefficient, but they did not provide explicit kernel circuits or a framework for deciding which radices are useful.

For general matrix polynomials, the Paterson–Stockmeyer (PS) method [10] achieves $O(\sqrt{n})$ products for degree-n evaluation, and Fasi and Higham [11] proved optimality within the PS scheme. Composition methods can do better on special polynomial families: Sastre and Ibáñez [12] achieved degree 15 in 4 products, exceeding the degree 9 reachable by PS with 4 products.

Our radix kernels import that composition viewpoint into the geometric-series setting. In particular, they extend Gustafsson et al.'s quinary construction [5] to higher radices, and the accompanying kernel-to-update framework shows how better kernel circuits translate directly into faster Neumann-series algorithms.

4. Radix-9 Kernel

The quinary kernel [5] achieves coefficient $4/\log_2 5 \approx 1.72$. The next natural target is radix 9. To improve on 1.72, a radix-9 update must use at most 5 products. Since $C(9) = \mu_9 + 2$, this means $\mu_9 \leq 3$. The degree-doubling bound already gives $\mu_9 \geq 3$, so a 3-product kernel would be optimal.

Derivation strategy. To hit this bound, we build the third product around $U = B^2$ and $V = U(B + 2U) = B^3 + 2B^4$, which are available after two products. We then seek one additional product $P \cdot Q$, with P and Q linear combinations of $\{B, U, V\}$, whose convolution fixes the high-degree terms $B^5, \ldots, B^8$. The remaining degrees B^2, B^3, B^4 can then be corrected using only additions of U and V. Matching the target coefficients gives an 8×8 linear system for the coefficients of P and Q.

Theorem 4.1 (Radix-9 kernel in 3 products)**.** *The kernel $T_9(B) = I + B + \cdots + B^8$ is computed by:*

1. $U := B^2$ *(MM1)*
2. $V := U(B + 2U) = B^3 + 2B^4$ *(MM2)*
3. $P := \frac{3}{20}B + 2U + V, \quad Q := \frac{11}{40}B - \frac{1}{8}U + \frac{1}{4}V$
4. $W := P \cdot Q$ *(MM3)*
5. $T_9(B) := W + I + B + \frac{767}{800}U + \frac{15}{32}V$

The key insight is that one carefully chosen product $P \cdot Q$ simultaneously generates four high-degree terms (B^5 through B^8), while the low-degree corrections adjust B^2–B^4 using only additions.

Proof. Let $\mathbf{p} = [\frac{3}{20}, 2, 1, 2]$ and $\mathbf{q} = [\frac{11}{40}, -\frac{1}{8}, \frac{1}{4}, \frac{1}{2}]$ be the coefficients of P and Q in powers B^1, B^2, B^3, B^4. Expanding $P \cdot Q$ shows that the coefficients of B^5 through B^8 are exactly 1. The same expansion gives coefficients $\frac{33}{800}$, $\frac{17}{32}$, and $\frac{1}{16}$ on B^2, B^3, and B^4. Adding $\frac{767}{800}U + \frac{15}{32}V$ raises those coefficients to 1, so the resulting polynomial is $I + B + \cdots + B^8$. □

Remark 4.2 (Rational coefficients). The radix-9 kernel uses only rational coefficients (fractions with denominators dividing 800), enabling exact computation in rational or fixed-point arithmetic. This contrasts with radix-15 (Section 5), where numerical optimization yields only real

coefficients with nonzero "spillover" into higher-degree terms. The transition from exact rational solutions (radix-5, radix-9) to approximate real solutions (radix-15) suggests that any exact higher-radix construction, if it exists, is substantially harder to uncover.

Having established the 3-product construction, we now quantify its benefit. With $\mu_9 = 3$, the update costs $C(9) = 5$ products. Iterating $t = \log_9 k$ times gives $S_k(A)$ in $5 \log_9 k \approx 1.58 \log_2 k$ products—a 21% improvement over binary and 8% over quinary.

In summary, the radix-9 kernel achieves the best known asymptotic coefficient among exact constructions, at the cost of one additional matrix product per update. The next section explores whether higher radices can improve further.

5. Certified Approximate Higher-Radix Kernels

The radix-9 kernel is *exact* and lower-bound optimal. For higher radices, the same degree-counting argument still points to useful coefficients—$\mu_{15} \geq 4$ and $\mu_{24} \geq 5$—but it does not tell us whether exact kernels exist or how approximate ones should be used. This section therefore separates two logically distinct tasks: *numerical discovery* of promising circuits, and *certified existence* of nearby prefix-exact kernels. Keeping those tasks separate prevents an optimizer output from being mistaken for a theorem.

5.1 Search Ansatz

We use the bilinear chained-product ansatz already implicit in the exact radix-9 construction. Fix $P_1 = B^2$. For a K-product kernel, define basis sets

$$\mathcal{B}_1 = \{I, B, P_1\}, \qquad \mathcal{B}_{i+1} = \mathcal{B}_i \cup \{P_i\},$$

and for $i = 2, \ldots, K$ write

$$L_i = \sum_{X \in \mathcal{B}_i} \ell_{i,X} X, \qquad R_i = \sum_{X \in \mathcal{B}_i} r_{i,X} X, \tag{4}$$
$$P_i = L_i R_i.$$

The final kernel is

$$f(B) = I + \gamma_1 B + \gamma_2 P_1 + \cdots + \gamma_K P_{K-1} + P_K. \tag{5}$$

For radix 15 we use $K = 4$, giving 28 parameters. For radix 24 we use $K = 5$, giving 41 parameters.

5.2 Numerical Discovery

Let $[f(B)]_j$ denote the coefficient of B^j. The discovery objective is the prefix residual

$$\Phi_m(\theta) = \sum_{j=0}^{m-1} \left([f_\theta(B)]_j - 1\right)^2,$$

where θ collects all ansatz parameters. We minimize Φ_m by multi-start local optimization. This stage is deliberately heuristic: it identifies promising seeds but does not by itself justify theorem-grade claims.

For radix 15 the best 4-product seed drives the prefix error below 5×10^{-16} while leaving spillover at degrees 15 and 16. For radix 24 the rounded appendix coefficients are insufficiently accurate for certification, so we refine them numerically before freezing the final seed.

5.3 Certified Approximate Kernels

The next step resolves the main expectation created by the search: each promising seed is converted into a *certified* kernel by solving the prefix equations exactly in a nearby parameter box. Write

$$F_m(\theta) = ([f_\theta]_0 - 1, \ldots, [f_\theta]_{m-1} - 1).$$

At each seed θ_0 we compute $DF_m(\theta_0)$, extract a well-conditioned square active subsystem, and apply a Krawczyk inclusion test on a small box in those active variables while holding the remaining variables fixed.

Theorem 5.1 (Certified approximate radix-15 and radix-24 kernels)**.** *There exists a certified radix-15 kernel in the 4-product bilinear ansatz and a certified radix-24 kernel in the 5-product bilinear ansatz such that the prefix equations $F_{15}(\theta) = 0$ and $F_{24}(\theta) = 0$ hold exactly. The corresponding certified boxes have maximum active-coordinate radii below 4.5×10^{-14} and 1.5×10^{-10}, respectively. Evaluating the kernels and their error maps over those boxes yields interval enclosures for every displayed spillover coefficient.*

Proof. The certification script `certify_kernel.py` implements the active-set Krawczyk test and evaluates the resulting coefficient intervals. The machine-readable certificates are `certified_coeffs_radix15.json` and `certified_co` a concise summary is given in Appendix Appendix 4. □

For radix 15, the certified error map is centered at

$$\begin{aligned} E_{15}(z) = {} & 0.8147910024806893\, z^{15} \\ & - 0.2731510277211418\, z^{16} \\ & + 0.4583600252404543\, z^{17}, \end{aligned}$$

with interval radii below 2.1×10^{-13}. For radix 24, the certified error map is centered at

$$\begin{aligned} E_{24}(z) = {} & 0.3551702076534789\, z^{24} + 0.4309645579871833\, z^{25} \\ & + 0.1654049800790244\, z^{26} - 0.1394687687945098\, z^{27} \\ & - 0.2416800071277084\, z^{28} - 0.1015904946346608\, z^{29} \\ & + 0.1098721022155451\, z^{30} + 0.2068198204262950\, z^{31} \\ & + 0.1575634044917271\, z^{32} + 0.0569441977034457\, z^{33}, \end{aligned}$$

with interval radii below 3.7×10^{-9}. These are the authoritative coefficients used in Sections 7 and 8; the appendix tables remain rounded display versions.

5.4 Local Structure

Certification also yields local geometric information about the solution set, clarifying why approximate kernels are natural in these ansätze. For the certified radix-15 kernel, the prefix Jacobian has full row rank 15 in the 28-parameter ansatz, so the local prefix-exact set is a smooth manifold of dimension $28 - 15 = 13$. For radix 24, the corresponding dimension is $41 - 24 = 17$. Appending the leading spillover coefficient as one more constraint raises the rank by one in both cases, so the local dimension drops by one.

This is still only a *local* result: it explains why prefix exactness plus small spillover is a natural optimization target near the certified solutions, but it does not prove global nonexistence of exact higher-radix kernels. Appendix Appendix 5 records the corresponding Jacobian-rank statement.

The remainder of the paper treats these certified approximate kernels as mathematical objects rather than optimizer outputs. Section 6 shows how spillover is absorbed by residual iteration, and Section 7 turns the certified error maps into matrix-norm decay bounds.

6. General Radix-Kernel Framework

The best known exact kernel is radix-9, with coefficient $5/\log_2 9 \approx 1.58$. Higher radices can do better in principle. A 4-product radix-15 kernel would drop the coefficient to $6/\log_2 15 \approx 1.54$, and a 5-product radix-24 kernel to $7/\log_2 24 \approx 1.53$. Our certified higher-radix kernels reach those coefficients, but they do so with nonzero *spillover* at degrees ≥ 15, so the standard exact update no longer applies. The issue is therefore not whether higher-radix kernels help, but how to use approximate ones without giving up the asymptotic gain.

Spillover invalidates the telescoping identity $T_m(A^n)(I - A^n) = I - A^{mn}$ that normally lets us extract A^{mn}. With spillover the identity fails and the power cannot be recovered. Hensel lifting $(Y' = Y(2I - MY))$ sidesteps power extraction, but its coefficient is 2.00—no better than binary splitting. We need an iteration that keeps the 1.54–1.53 promise of higher-radix kernels without exact power extraction.

We resolve both obstacles by applying the kernel to the *residual* rather than to powers of A, eliminating the need for power extraction. We develop this in three stages: formalize approximate kernels, analyze their composition as polynomials, then lift to matrices.

6.1 Approximate Radix Kernels

We first make precise what "approximate kernel" means and how its error behaves.

Definition 6.1 (Approximate radix-m kernel)**.** A polynomial $f \in \mathbb{R}[z]$ is an *approximate radix-m kernel* if $(1 - z)f(z) = 1 + O(z^m)$. Equivalently, $f(z) \equiv (1-z)^{-1} \pmod{z^m}$. The terms of degree $\geq m$ in f are called *spillover*. When the spillover is zero, the kernel is *exact* and coincides with $T_m(z)$.

Definition 6.2 (Error map)**.** Given an approximate radix-m kernel f, define $E(z) = 1 - (1 - z)f(z)$.

Lemma 6.3 (Error order)**.** *If $f(z) = \sum a_j z^j$ is an approximate radix-m kernel, then*

$$E(z) = (1 - a_m)\, z^m + O(z^{m+1}).$$

In particular, $c = 1 - [f]_m$. If f is exact, $a_m = 0$, so $c = 1$ and $E(z) = z^m$.

Proof. Since $(1 - z)f(z) = 1 + O(z^m)$, the degree-m term is $(1 - z)f(z) = 1 + (a_m - a_{m-1})\, z^m + \cdots = 1 + (a_m - 1)\, z^m + \cdots$, so $E(z) = 1 - (1 - z)f(z) = (1 - a_m)\, z^m + O(z^{m+1})$. □

6.2 General Radix-Kernel Summation

The key idea is to apply the kernel to its own error: $f(E(z))$ corrects the residual left by $f(z)$.

Definition 6.4 (General radix-kernel summation)**.** Given an approximate radix-m kernel f with error map E, one *lift* of the general radix-kernel summation produces

$$f(z) \cdot f\big(1 - (1 - z)f(z)\big) = f(z) \cdot f(E(z)). \tag{6}$$

Lemma 6.5 (Composition identity)**.** *For any approximate radix-m kernel f with error map E:*

$$E(E(z)) = 1 - (1 - z)f(z)f(E(z)). \tag{7}$$

Proof. By definition, $1 - E(z) = (1 - z)f(z)$. Applying the definition of E at the point $E(z)$:

$$E(E(z)) = 1 - (1 - E(z))f(E(z)) = 1 - (1 - z)f(z)f(E(z)).$$

□

Lemma 6.6 (Prefix growth)**.** *If f is an approximate radix-m kernel, then $f(z) \cdot f(E(z))$ matches the first m^2 coefficients of $(1 - z)^{-1}$.*

Proof. By Lemma 6.5, $(1-z)f(z)f(E(z)) = 1-E(E(z))$. By Lemma 6.3, $E(z) = O(z^m)$, so $E(E(z)) = O(z^{m^2})$. Hence $(1 - z)f(z)f(E(z)) = 1 - O(z^{m^2})$, and $f(z)f(E(z))$ agrees with $(1 - z)^{-1}$ to order $m^2 - 1$. □

Lemma 6.7 (Error propagation). *Write $E^{[n]}$ for the n-fold composition ($E^{[0]}(z) = z$, $E^{[n+1]} = E \circ E^{[n]}$). Let $E(z) = c\,z^m + O(z^{m+1})$. Then*

$$E^{[n]}(z) = c^{(m^n-1)/(m-1)}\, z^{m^n} + O(z^{m^n+1}).$$

Spillover is harmless to the prefix: regardless of c, it is pushed to degree $\geq m^n$ after n steps.

Proof. $E(E(z)) = c\,(cz^m)^m + \cdots = c^{m+1}\, z^{m^2} + \cdots$. By induction, $E^{[n]}(z) = c^{1+m+\cdots+m^{n-1}}\, z^{m^n} + \cdots = c^{(m^n-1)/(m-1)}\, z^{m^n} + \cdots$. □

Remark 6.8 (Analytic stability). In floating-point arithmetic, the coefficient magnitude matters. When $|c| < 1$, the leading coefficient $c^{(m^n-1)/(m-1)}$ decays, bounding the residual norm. For the radix-15 kernel (Section 5), $[f]_{15} \approx 0.19$ gives $c \approx 0.81 < 1$.

6.3 Matrix Iteration

We now translate this polynomial machinery to matrices. Let $M = I - A$. We maintain an approximate inverse $Y_n \approx M^{-1}$ and its residual $R_n = I - MY_n$. Starting from $Y_0 = I$ and $R_0 = A$, each step applies f to the residual:

$$Y_{n+1} = Y_n \cdot f(R_n), \qquad R_{n+1} = I - M\,Y_{n+1}. \tag{8}$$

Lemma 6.9 (General radix-kernel iteration). *Under iteration (8):*

1. $R_n = E^{[n]}(A)$;
2. $R_n = O(A^{m^n})$;
3. $Y_n = \sum_{j=0}^{m^n-1} A^j + O(A^{m^n})$.

Proof. Base case. $R_0 = A = E^{[0]}(A)$. *Induction.* If $R_n = E^{[n]}(A)$, then $R_{n+1} = I - M\,Y_n f(R_n) = I - (I - R_n) f(R_n) = E(R_n) = E^{[n+1]}(A)$. Parts (2)–(3) follow from $E(z) = O(z^m)$ and $M\,Y_n = I - R_n = I - O(A^{m^n})$. □

6.4 Cost Analysis

We now quantify that gain and compare against prior methods. Each step (8) requires μ_m products for $f(R_n)$, 1 for $Y_{n+1} = Y_n \cdot f(R_n)$, and 1 for $R_{n+1} = I - M\,Y_{n+1}$: total $\mu_m + 2$ per step—identical to the exact framework. Achieving accuracy $k = m^t$ requires $t = \log_m k$ steps:

$$(\mu_m + 2)\log_m k = \frac{\mu_m + 2}{\log_2 m}\,\log_2 k. \tag{9}$$

Corollary 6.10 (Radix-15 and radix-24 coefficients). *An approximate radix-15 kernel with $\mu_{15} = 4$ achieves coefficient $6/\log_2 15 \approx 1.54$. An approximate radix-24 kernel with $\mu_{24} = 5$ achieves coefficient $7/\log_2 24 \approx 1.53$.*

Table 1. Matrix product counts for $S_k(A)$ with $k \geq 10^p$. Cost per step is $\mu_m + 2$, minus 2 for the trivial first update ($Y_0 = I$) and the unnecessary last residual. Radix-24$^+$ uses radix-24 for all steps except possibly the last, where a smaller radix may be cheaper.

Method	Coeff.	10^3	10^6	10^9	10^{12}	10^{15}
Binary [6]	2.00	18	38	58	78	98
Quinary [5]	1.72	18	34	50	70	86
Radix-9	1.58	18	33	48	63	78
Radix-15	1.54	16	34	46	64	76
Radix-24	1.53	19	33	47	61	75
Radix-24$^+$	—	**14**	**29**	**44**	**59**	**75**

The certified radix-15 and radix-24 kernels satisfy the prefix equations exactly at the parameter level, so the coefficient model applies to the mathematical iteration itself. Section 8 then checks that floating-point implementations exhibit the predicted m^n prefix growth and the expected product-count savings.

Table 1 compares all methods against prior work. Under naive kernel computation, higher radices were counterproductive and left coefficients at or above $2\log_2 k$. Once low-product kernels are available, the ordering reverses and higher radices become the asymptotically best options.

Since any step can use any available kernel, the optimal strategy uses radix-24 for all but the final step, where a smaller radix may suffice: e.g. for $k \geq 10^3$, the sequence $(24, 24, 2)$ gives $k = 1152$ in just 14 products. At $k = 10^{15}$, this saves 23 products over binary (75 vs 98, a 23% reduction).

The optimization principle: choose m to minimize $C(m)/\log_2 m$. For our kernels: $4/\log_2 5 \approx 1.72$ (quinary), $5/\log_2 9 \approx 1.58$ (radix-9), $6/\log_2 15 \approx 1.54$ (radix-15), $7/\log_2 24 \approx 1.53$ (radix-24). Higher radices help only if kernel cost grows sublinearly in $\log_2 m$.

6.5 Special Cases

The framework recovers known methods as special cases.

Hensel lifting ($m = 2$). $f(z) = 1 + z$, $E(z) = z^2$, $\mu_2 = 0$: $Y_{n+1} = Y_n(I + R_n) = Y_n(2I - MY_n)$, coefficient $2/\log_2 2 = 2.00$.

Exact radix-m. $f = T_m$, $E(z) = z^m$, $R_n = A^{m^n}$: recovers the standard telescoping update.

The framework resolves both obstacles by applying the kernel to the residual, not to powers of A. The cost per step is $\mu_m + 2$—identical to the exact case—and the prefix grows as m^n regardless of spillover. Hensel lifting and exact telescoping emerge as special cases, so the framework subsumes all prior methods. With radix-24 it achieves

coefficient 1.53, the best known rate. The cost analysis assumes convergence; the next section examines when spillover preserves it.

7. Stability Analysis

Section 6 showed that the residual satisfies $R_n = E^{[n]}(A)$, where

$$E(z) = 1 - (1-z)f(z)$$

is the kernel error map. Prefix growth alone does not imply bounded iterates: once spillover is present, the relevant question is whether repeated composition of E drives residuals to zero on the matrix set of interest. This section makes that question quantitative.

7.1 Error-Map Coefficients

Let

$$f(z) = \sum_{j=0}^{d} a_j z^j, \qquad a_0 = \cdots = a_{m-1} = 1,$$

so that the prefix is exact through degree $m-1$. Then

$$E(z) = \sum_{j=m}^{d+1} e_j z^j, \tag{10}$$

with

$$e_m = 1 - a_m, \qquad e_j = a_{j-1} - a_j \ (m < j \le d), \qquad e_{d+1} = a_d.$$

The leading coefficient e_m is only the first spillover datum; all of (10) matters for stability.

7.2 Scalar Tests

Define

$$q_\rho := \frac{1}{\rho} \max_{|z|=\rho} |E(z)| = \frac{1}{\rho} \max_{|z|\le\rho} |E(z)|, \qquad \eta_\rho := \sum_{j=m}^{d+1} |e_j| \rho^{j-1},$$

and

$$s_0 := 0, \qquad s_n := 1 + m + \cdots + m^{n-1} = \frac{m^n - 1}{m-1} \quad (n \ge 1).$$

Basin test. For a scalar state $\lambda_{n+1} = E(\lambda_n)$, convergence requires $E^{[n]}(\lambda) \to 0$. This is the exact scalar criterion.

Disk test. If $q_\rho \le 1$, then the disk $|z| \le \rho$ is forward invariant. Since $E(z) = z^m h(z)$ has a zero of order m at the origin, the generalized Schwarz lemma gives

$$|E^{[n]}(z)| \le \rho\, q_\rho^{s_n} \Big(\frac{|z|}{\rho}\Big)^{m^n}, \qquad |z| \le \rho. \tag{11}$$

Coefficient surrogate. The cheaper quantity η_ρ satisfies $q_\rho \le \eta_\rho$. Hence $\eta_\rho \le 1$ is a sufficient, though possibly conservative, way to verify the disk test.

7.3 A Functional-Calculus Theorem

The scalar bound (11) lifts immediately to any matrix class that admits a polynomial functional-calculus estimate.

Theorem 7.1 (Abstract matrix bound). *Let $\Omega \subset \mathbb{C}$ be compact, and suppose*

$$\sigma(A) \subset \Omega \subset \{z : |z| \le \rho\}, \qquad \|p(A)\|_2 \le M_\Omega \max_{z\in\Omega} |p(z)|$$

for every polynomial p. Set $r_\Omega := \max_{z\in\Omega} |z|$. If $q_\rho \le 1$, then for every $n \ge 0$,

$$\|R_n\|_2 \le M_\Omega\, \rho\, q_\rho^{s_n} \Big(\frac{r_\Omega}{\rho}\Big)^{m^n}.$$

Consequently,

$$\frac{\|(I-A)^{-1} - Y_n\|_2}{\|(I-A)^{-1}\|_2} \le M_\Omega\, \rho\, q_\rho^{s_n} \Big(\frac{r_\Omega}{\rho}\Big)^{m^n}.$$

Proof. Lemma 6.9 gives $R_n = E^{[n]}(A)$. Applying the assumed functional-calculus estimate to the polynomial $E^{[n]}$ and then using (11) on Ω yields the residual bound. Since $(I-A)^{-1} - Y_n = (I-A)^{-1} R_n$, the relative inverse-error estimate follows immediately. □

Corollary 7.2 (Normal matrices). *If A is normal, then Theorem 7.1 applies with $\Omega = \sigma(A)$ and $M_\Omega = 1$. In particular,*

$$\|R_n\|_2 \le \rho\, q_\rho^{s_n} \Big(\frac{\rho(A)}{\rho}\Big)^{m^n},$$

whenever $\rho \ge \rho(A)$ and $q_\rho \le 1$.

Corollary 7.3 (Numerical-range bound). *Let*

$$W(A) = \{x^* A x : \|x\|_2 = 1\}, \qquad w(A) = \max_{z\in W(A)} |z|.$$

By Crouzeix and Palencia [13],

$$\|p(A)\|_2 \le C_{\mathrm{Cr}} \max_{z\in W(A)} |p(z)|, \qquad C_{\mathrm{Cr}} \le 1 + \sqrt{2}.$$

Hence, if $\rho \ge w(A)$ and $q_\rho \le 1$, then

$$\|R_n\|_2 \le C_{\mathrm{Cr}}\, \rho\, q_\rho^{s_n} \Big(\frac{w(A)}{\rho}\Big)^{m^n}.$$

If moreover $\rho < 1$, then

$$\|(I-A)^{-1} - Y_n\|_2 \le \frac{C_{\mathrm{Cr}}}{1-\rho}\, \rho\, q_\rho^{s_n} \Big(\frac{w(A)}{\rho}\Big)^{m^n}.$$

Proposition 7.4 (Pseudospectral refinement). *Let Γ_ε be a rectifiable contour enclosing $\sigma(A)$ such that*

$$\|(zI - A)^{-1}\|_2 \le \varepsilon^{-1} \qquad \text{for all } z \in \Gamma_\varepsilon,$$

and let $L_\varepsilon = \mathrm{len}(\Gamma_\varepsilon)$. Then for every polynomial p,

$$\|p(A)\|_2 \le \frac{L_\varepsilon}{2\pi\varepsilon} \max_{z\in\Gamma_\varepsilon} |p(z)|.$$

If $\Gamma_\varepsilon \subset \{z : |z| \le \rho\}$ and $r_\varepsilon := \max_{z\in\Gamma_\varepsilon} |z|$, then

$$\|R_n\|_2 \le \frac{L_\varepsilon}{2\pi\varepsilon}\, \rho\, q_\rho^{s_n} \left(\frac{r_\varepsilon}{\rho}\right)^{m^n}.$$

Proof. The contour bound is the standard polynomial functional calculus on Γ_ε; see, for example, [3, Chapter 17] or [14]. The final estimate is another application of Theorem 7.1. □

7.4 Certified Radii for Radix 15

For the certified radix-15 kernel, Appendix Appendix 4 centers the error map at

$$\begin{aligned} E_{15}(z) = {} & 0.8147910024806893\, z^{15} \\ & - 0.2731510277211418\, z^{16} \\ & + 0.4583600252404543\, z^{17}. \end{aligned}$$

Here the coefficient surrogate is essentially sharp:

$$\begin{aligned} \eta_{15}(\rho) = {} & 0.8147910024806893\, \rho^{14} \\ & + 0.2731510277211418\, \rho^{15} \\ & + 0.4583600252404543\, \rho^{16}. \end{aligned}$$

The certified coefficient radius is

$$\rho_{\eta,15} = 0.970897959\ldots,$$

and the certified direct numerical-range radius is

$$\rho_{q,15}^{\mathrm{cert}} = 0.970699488\ldots.$$

At the midpoint coefficients, direct circle maximization gives the estimated value

$$\rho_{q,15}^{\mathrm{est}} = 0.970897959\ldots,$$

which coincides with $\rho_{\eta,15}$ to the displayed digits. Thus, for practical purposes, the coefficient test is sharp for radix 15.

At $w(A) = \rho = 0.95$, Corollary 7.3 yields

$$\|R_n\|_2 \le C_{\mathrm{Cr}}\, 0.95\, (0.725638)^{(15^n-1)/14},$$

so

$$\|R_2\|_2 \lesssim 1.36 \times 10^{-2}, \qquad \|R_3\|_2 \lesssim 6.3 \times 10^{-34}.$$

7.5 Certified Radii for Radix 24

For the certified radix-24 kernel, the error map is centered at

$$\begin{aligned} E_{24}(z) = {} & 0.3551702076534789\, z^{24} + 0.4309645579871833\, z^{25} + 0.165 \\ & - 0.1394687687945098\, z^{27} - 0.2416800071277084\, z^{28} - 0.1 \\ & + 0.1098721022155451\, z^{30} + 0.2068198204262950\, z^{31} + 0.1 \\ & + 0.0569441977034457\, z^{33}. \end{aligned}$$

The coefficient surrogate certifies

$$\rho_{\eta,24} = 0.974559738\ldots,$$

while the interval-based direct circle test certifies the larger radius

$$\rho_{q,24}^{\mathrm{cert}} = 0.984176086\ldots.$$

At the midpoint coefficients, direct circle maximization gives the estimated value

$$\rho_{q,24}^{\mathrm{est}} = 0.984432110\ldots.$$

The gap between $\rho_{\eta,24}$ and $\rho_{q,24}^{\mathrm{cert}}$ is already visible in the certified intervals, so cancellations matter substantially for radix 24.

7.6 Summary

The error-map viewpoint separates three levels of statement:

- **exact scalar criterion:** $E^{[n]}(\lambda) \to 0$ for each scalar state of interest;
- **certified scalar tests:** $q_\rho \le 1$ or the cheaper $\eta_\rho \le 1$;
- **matrix-level corollaries:** normal, numerical-range, and pseudospectral bounds through Theorem 7.1.

For radix 15, the coefficient test is effectively sharp and certifies numerical-range radii near 0.971. For radix 24, the direct circle test is materially better, certifying radii near 0.984 while the coefficient surrogate stops near 0.975. Section 8 compares these theorem predictions with observed matrix residuals.

8. Numerical Experiments

The experiments now play two roles. First, they validate the theorem-driven stability picture from Section 7: the certified radii, the gap between η_ρ and q_ρ, and the effect of nonnormality on matrix bounds. Second, they confirm that the certified higher-radix kernels still deliver the predicted product-count savings in practical inverse approximation.

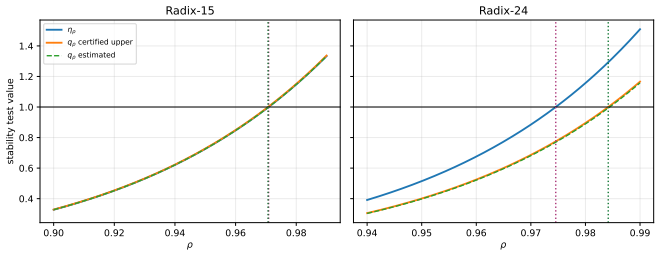

Figure 1. Certified scalar stability curves for the radix-15 and radix-24 kernels. The plotted data are stored in `qrho_curves.csv`, and the certified crossing radii are summarized in `qrho_radii.csv`.

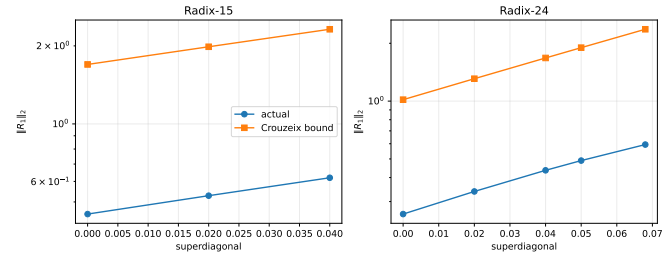

Figure 2. Observed residual norms versus the numerical-range bound on normal and nonnormal test families. The bound remains valid and becomes looser as nonnormality increases, as expected from the Crouzeix factor.

Reproducibility. All tables and theorem-validation figures in this section are generated by `certify_kernel.py` and `reproduce_theory_figures.py`. The mapping from output artifact to script is recorded in the supplement manifest `reproducibility_manifest.txt`. All computations use Python with NumPy/SciPy and IEEE 754 double precision.

8.1 Certified Radius Curves

Figure 1 plots the cheap coefficient surrogate η_ρ, the certified upper bound for q_ρ, and midpoint circle maximization for the certified radix-15 and radix-24 kernels. The picture matches the theory from Section 7: for radix 15 the coefficient surrogate is essentially sharp, whereas for radix 24 the direct circle test is noticeably better.

The certified radii are

$$\rho_{\eta,15} = 0.970897959, \qquad \rho_{q,15}^{\mathrm{cert}} = 0.970699488,$$

and

$$\rho_{\eta,24} = 0.974559738, \qquad \rho_{q,24}^{\mathrm{cert}} = 0.984176086.$$

For radix 24, midpoint circle maximization raises the estimated crossing to 0.984432110, so the certified direct test captures most of the available cancellation already.

8.2 Normal and Nonnormal Validation

Figure 2 compares actual residual norms with the numerical-range bound from Corollary 7.3 on two small families. The first family is normal, with diagonal matrices at radii just inside the certified threshold. The second family is nonnormal, using 2×2 upper-triangular matrices with fixed diagonal and increasing superdiagonal. The corresponding raw data appear in `theory_validation.csv`.

Two qualitative points stand out. First, the certified bounds are conservative but correctly ordered: the actual residual decays far below the bound in the normal family, then moves closer to it as the matrix becomes more nonnormal. Second, the radius tests are not cosmetic. The failure-side scalar traces in `failure_side.csv` show that crossing the certified threshold materially worsens the one-step residual even when later compositions may still decay for a particular scalar point.

Table 2. Product counts for $S_k(A)$ at pure powers (binary baseline).

k	Binary	Quinary	Radix-9	Radix-15	Radix-24	Savings
$125 = 5^3$	14	**12**	—	—	—	14%
$625 = 5^4$	20	**16**	—	—	—	20%
$81 = 9^2$	14	—	**10**	—	—	29%
$729 = 9^3$	20	—	**15**	—	—	25%
$225 = 15^2$	16	—	—	**12**	—	25%
$3375 = 15^3$	24	—	—	**18**	—	25%
$576 = 24^2$	20	—	—	—	**14**	30%
$13824 = 24^3$	28	—	—	—	**21**	25%

8.3 Product Counts and Approximate Inversion

Stability certification matters only if the faster kernels still reduce work, so we close by returning to the original efficiency question. Table 2 lists the exact matrix-product counts for pure powers. The higher-radix certified kernels still deliver the same asymptotic improvements as in the original draft: radix 9 reduces the binary constant from 2.00 to 1.58, radix 15 reaches 1.54, and radix 24 reaches 1.53.

Figure 3 keeps the original convergence comparison on a normal test family, while Figure 4 records the asymptotic coefficients themselves. These are now interpreted as a secondary layer: once the certified kernels are available and stable on the relevant matrix set, the higher radix translates directly into fewer products and lower runtime.

Summary. The experiments now line up with the theory in three layers. The scalar curves explain which radii are certified. The normal and nonnormal matrix tests validate the matrix bounds. The remaining convergence and product-count plots show that once those stability con-

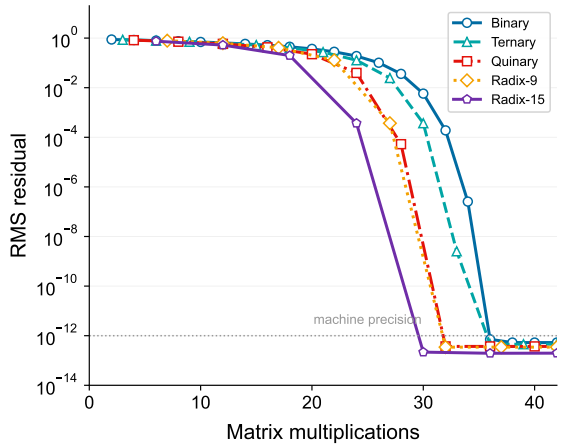

Figure 3. Convergence comparison on a normal family. Higher-radix methods reach machine precision with fewer matrix products once the certified stability conditions hold.

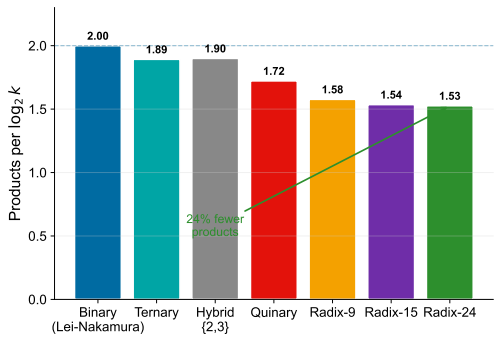

Figure 4. Asymptotic coefficients (matrix products per $\log_2 k$). The certified radix-24 kernel retains the best asymptotic rate at 1.53.

ditions are met, the certified higher-radix kernels realize the expected 25–30% reduction in matrix products.

9. Conclusion

We set out to extend Neumann series evaluation beyond the known radix-5 kernel [5], where no higher-radix constructions existed and where nonexact higher-radix kernels lacked a matrix-level convergence theory. We addressed this gap with four linked contributions:

1. The radix-9 kernel achieves $\mu_9 = 3$ with rational coefficients. It is the first exact construction beyond radix 5 and yields $5 \log_9 k \approx 1.58 \log_2 k$ products, 21% fewer than binary splitting.
2. Certified approximate radix-15 and radix-24 kernels, obtained by numerical discovery plus interval certification, yield $6/\log_2 15 \approx 1.54$ and $7/\log_2 24 \approx 1.53$.
3. A general radix-kernel framework (Section 6) applies kernels to residuals, enabling iterative use of approximate kernels despite spillover.
4. Matrix-level convergence follows from an abstract theorem with normal, numerical-range, and pseudospectral corollaries that quantify residual decay for nonnormal matrices.

The transition from exact (radix-9) to certified approximate (radix-15, radix-24) kernels suggests that spillover is the norm at higher radices. Our local structure analysis supports that interpretation without claiming a global impossibility theorem: near the certified solutions, the prefix-exact set is a positive-dimensional manifold and the leading-spillover cancellation imposes one additional local constraint.

The reproducibility workflow generates the certified radii, theorem-validation figures, and machine-readable coefficient files from one stack of scripts and supplementary artifacts. Those materials make the paper's claims auditable, and the public companion repository and tutorial site are staged for release at the project GitHub repository and the tutorial site.

This leaves several mathematical directions:

1. *Exact vs. approximate boundary*: Can we prove that exact radix-15 kernels do not exist, or find one?
2. *Higher radices*: Can approximate kernels with μ_m growing sublogarithmically improve the coefficient beyond 1.53?
3. *Structured matrices*: How do these algorithms extend to sparse or Toeplitz settings?

More broadly, the paper reframes truncated Neumann-series evaluation as a kernel-design problem: better kernels translate directly into better series algorithms. That viewpoint may also benefit log-determinant estimation [4], where polynomial traces from radix-kernel sums could replace individual power traces while reducing the number of matrix products.

Acknowledgements

This work was supported by the Laboratory Directed Research and Development Program of Oak Ridge National Laboratory, managed by UT-Battelle, LLC, for the U.S. Department of Energy.

References

[1] Paul F. Dubois, Anne Greenbaum, and Garry H. Rodrigue. "Approximating the Inverse of a Matrix for Use in Iterative Algorithms on Vector Processors". In: *Computing* 22.3 (1979), pp. 257–268. DOI: 10.1007/BF02243566.

[2] Yousef Saad. *Iterative Methods for Sparse Linear Systems*. 2nd. SIAM, 2003. DOI: 10.1137/1.9780898718003.

[3] Nicholas J. Higham. *Functions of Matrices: Theory and Computation*. SIAM, 2008. DOI: 10.1137/1.9780898717778.

[4] Piyush Sao. "What Trace Powers Reveal About Log-Determinants: Closed-Form Estimators, Certificates, and Failure Modes". In: *arXiv preprint arXiv:2601.12612* (2025). URL: https://arxiv.org/abs/2601.12612.

[5] Oscar Gustafsson et al. "Approximate Neumann Series or Exact Matrix Inversion for Massive MIMO?" In: *Proc. IEEE Symposium on Computer Arithmetic (ARITH)*. 2017, pp. 62–63. DOI: 10.1109/ARITH.2017.14.

[6] Shao-Feng Lei and Tadao Nakamura. "On the Fast Evaluation of Parallel ARMA Filters and Related Matrix Power Sums". In: *IEEE Transactions on Circuits and Systems I* 39.4 (1992), pp. 299–300. DOI: 10.1109/81.129460.

[7] Vassil S. Dimitrov and Todor V. Cooklev. "Two Algorithms for Modular Exponentiation Using Nonstandard Arithmetics". In: *IEICE Transactions on Fundamentals* E78-A.1 (1995), pp. 82–87.

[8] David Westreich. "A Practical Method for Computing the Exponential of a Matrix and its Integral". In: *IEEE Transactions on Circuits and Systems* 36.1 (1989), pp. 162–164. DOI: 10.1109/31.16585.

[9] Yusuke Matsumoto, Naofumi Takagi, and Kazuyoshi Takagi. "Efficient Algorithms for Computing Power Sums Using Higher Radices". In: *IEICE Transactions on Fundamentals* E101-A.2 (2018), pp. 467–471. DOI: 10.1587/transfun.E101.A.467.

[10] Michael S. Paterson and Larry J. Stockmeyer. "On the Number of Nonscalar Multiplications Necessary to Evaluate Polynomials". In: *SIAM Journal on Computing* 2.1 (1973), pp. 60–66. DOI: 10.1137/0202007.

[11] Massimiliano Fasi. "Optimality of the Paterson–Stockmeyer Method for Evaluating Matrix Polynomials and Rational Matrix Functions". In: *Linear Algebra and its Applications* 574 (2019), pp. 182–200. DOI: 10.1016/j.laa.2019.04.001.

[12] Jorge Sastre and Javier Ibáñez. "Boosting the Computation of the Matrix Exponential". In: *Mathematics* 9.14 (2021), p. 1600. DOI: 10.3390/math9141600.

[13] Michel Crouzeix and Carlos Palencia. "The Numerical Range Is a $(1+\sqrt{2})$-Spectral Set". In: *SIAM Journal on Matrix Analysis and Applications* 38.2 (2017), pp. 649–655. DOI: 10.1137/17M1116672.

[14] Lloyd N. Trefethen and Mark Embree. *Spectra and Pseudospectra: The Behavior of Nonnormal Matrices and Operators*. Princeton University Press, 2005.

Appendix 1. Kernel Coefficients

This appendix provides display formulas for the kernels and supplementary numerical validation on test matrices. For the certified approximate kernels, the authoritative machine-readable data are the supplementary JSON files described in Appendix Appendix 4.

Appendix 1.1 Quinary Kernel (2 Products)

The quinary kernel $T_5(B) = I + B + B^2 + B^3 + B^4$ is computed using 2 matrix products:

$$U = B^2, \tag{12}$$
$$V = U(B + U) = B^3 + B^4, \tag{13}$$
$$T_5(B) = I + B + U + V. \tag{14}$$

This is exact: all 5 coefficients equal 1.

Appendix 1.2 Radix-9 Kernel (3 Products)

The radix-9 kernel $T_9(B) = \sum_{j=0}^{8} B^j$ uses 3 products:

$$U = B^2, \tag{15}$$
$$V = U(B + 2U) = B^3 + 2B^4, \tag{16}$$
$$P = (0.15B + 2U + V)(0.275B - 0.125U + 0.25V), \tag{17}$$
$$T_9(B) = I + B + \tfrac{767}{800}U + \tfrac{15}{32}V + P. \tag{18}$$

The rational coefficients in (18) are $\frac{767}{800} = 0.95875$ and $\frac{15}{32} = 0.46875$.

Appendix 1.3 Numerical Validation

We validate convergence on a 64×64 symmetric positive definite (SPD) matrix A with geometrically spaced eigenvalues in $[10^{-2}, 1]$, giving $\kappa(I - A) \approx 10^4$. In this appendix, let $B = I - A$ (distinct from the kernel argument B in the main text).

Setup. For series index k, each method computes an approximation S_k to $(I - B)^{-1}$. We measure the relative residual $\|I - (I - B)S_k\|_F/\sqrt{d}$, where $\|\cdot\|_F$ denotes the Frobenius norm.

Results. Table 3 lists the matrix product count required to reach machine precision ($\approx 10^{-13}$).

Radix-24 uses the fewest products and, with refined coefficients, reaches machine precision despite nonzero spillover. The naive method requires $k - 1 = 10^4$ products for comparable accuracy.

Table 3. Matrix product counts to reach 10^{-13} residual ($d = 64$, $\kappa \approx 10^4$).

Method	Products	Residual
Binary ($m = 2$)	38	5.6×10^{-13}
Ternary ($m = 3$)	36	4.7×10^{-13}
Quinary ($m = 5$)	32	3.9×10^{-13}
Radix-9	32	4.1×10^{-13}
Radix-15 (approx.)	30	4.2×10^{-13}
Radix-24 (approx.)	28	3.4×10^{-14}

Efficiency comparison. Per radix step, the update cost is $C(m) = \mu_m + 2$ products:

- Binary: $C(2) = 0 + 2 = 2$, coefficient $\frac{2}{\log_2 2} = 2.00$;
- Ternary: $C(3) = 1 + 2 = 3$, coefficient $\frac{3}{\log_2 3} \approx 1.89$;
- Quinary: $C(5) = 2 + 2 = 4$, coefficient $\frac{4}{\log_2 5} \approx 1.72$;
- Radix-9: $C(9) = 3 + 2 = 5$, coefficient $\frac{5}{\log_2 9} \approx 1.58$;
- Radix-15: $C(15) = 4 + 2 = 6$, coefficient $\frac{6}{\log_2 15} \approx 1.54$;
- Radix-24: $C(24) = 5 + 2 = 7$, coefficient $\frac{7}{\log_2 24} \approx 1.53$.

Radix-9 achieves the best rate among exact kernels; radix-24 achieves the best overall.

Appendix 1.4 Approximate Kernels: Residual vs Radix Trade-off

For a single truncated Neumann computation (no recursion), we compare approximate kernels at fixed matrix product budgets (6 and 7 products). Test matrix: $d = 64$, geometrically spaced eigenvalues in $[0.01, 1]$, giving $\rho(B) = 0.99$.

Key finding: Higher radix dominates. Table 4 compares kernels at 6 products and 7 products budgets. Even with high kernel residual, higher radix gives better inverse approximation. Error is measured in spectral norm: $\|S - A^{-1}\|_2/\|A^{-1}\|_2$, where $\|\cdot\|_2$ denotes the largest singular value.

Interpretation. The "vs Exact" column shows degradation from kernel approximation error. Even radix-33 with residual 0.6 is only 0.7% worse than an exact radix-33 kernel. The dominant factor is radix m: error scales as $\rho(B)^m$, so larger m gives exponentially better approximation, overwhelming the kernel residual effect.

Coefficients for radix-33 kernel (5 products). The approximate coefficients c_j for $\widetilde{T}_{33}(B) = \sum_{j=0}^{32} c_j B^j$: Coefficients range from 0.69 to 1.27, yet the inverse error is only 1% worse than exact. No spillover terms exist beyond degree 32.

Table 4. Inverse approximation error (spectral norm) for single kernel application. "vs Exact": ratio to exact Neumann at same radix.

Method	Prods	Kern. Res.	Inv. Err.	vs Exact
6-product budget				
Standard S_7 (radix 7)	–	0	0.932	–
Radix-15	4	5×10^{-16}	0.860	1.000
Radix-16	4	1×10^{-2}	0.852	1.000
Radix-17	4	1×10^{-1}	0.844	1.001
7-product budget				
Standard S_8 (radix 8)	–	0	0.923	–
Radix-24	5	2×10^{-15}	0.786	1.000
Radix-27	5	1×10^{-3}	0.762	1.000
Radix-31	5	1×10^{-1}	0.733	1.002
Radix-33	5	6×10^{-1}	0.722	1.007

j	c_j values
0–10	1.00, 1.00, 1.00, 0.98, 1.00, 1.02, 1.00, 0.69, 1.12, 0.99, 1.05
11–21	0.88, 1.11, 0.69, 0.73, 1.01, 1.27, 1.04, 1.08, 1.06, 0.98, 1.00
22–32	0.99, 0.88, 1.13, 0.95, 1.02, 0.97, 1.01, 0.99, 1.00, 1.00, 1.00

Appendix 2. Radix-15 Kernel: Explicit Circuit

The circuit (5) with $P_1 = B^2$ is fully specified by the following coefficients, obtained by multi-stage refinement (L-BFGS-B, Nelder–Mead, Powell) from 200 random starts. Each product $P_i = L_i \cdot R_i$; rows of the table give coefficients of $\{I, B, P_1, \ldots\}$. The table is rounded for display; the authoritative certified coefficient and error-map intervals are stored in `certified_coeffs_radix15.json`.

	I	B	P_1	P_2	P_3
L_2	0.237957446781864	0.241306784696238	1.573677404167023		
R_2	0.047776286323906	−0.046512791269008	0.888968295817947		
L_3	0.262503282946680	0.918866331886906	0.841863255429817	0.819355416606328	
R_3	0.183860759371110	0.068445979922426	0.133854440389316	−1.405242484443339	
L_4	−0.004559362003080	−1.384977065724292	0.021866758593494	0.122021888194117	0.274863153197116
R_4	−0.701388642425678	0.601773558132688	−0.624151559177423	−0.137295755918682	0.328424149078130

Final combination:

$$\widetilde{T}_{15} = I + 0.077578347530026\,B + 1.778433959744106\,P_1 + 0.663737059861237\,P_2 - 0.024081230623538\,P_3 + P_4.$$

At these coefficients, the prefix is exact to double precision: $\max_{0\le j\le 14} |[\widetilde{T}_{15}]_j - 1| < 5\times10^{-16}$. The spillover is $0.185\,B^{15} + 0.458\,B^{16} + O(B^{17})$, giving error map leading coefficient $c = 1 - [\widetilde{T}_{15}]_{15} \approx 0.81$.

Appendix 3. Radix-24 Kernel: Explicit Circuit

The 5-product circuit with $P_1 = B^2$ is specified by the coefficients below, obtained by the same refinement procedure (seed 196 from 300 random starts). Each product $P_i = L_i \cdot R_i$. These coefficients are rounded display values. For certification and theorem generation we use the refined machine-readable seed recorded in `certified_coeffs_radix24.json`.

	I	B	P_1	P_2	P_3	P_4
L_2	0.328288	0.333461	−1.641330			
R_2	0.101660	0.635274	0.942490			
L_3	−0.403911	−0.185364	−0.019310	0.764381		
R_3	−0.683152	0.059370	0.955595	−0.835290		
L_4	0.850396	−0.170005	0.400796	−0.758391	−0.420233	
R_4	0.332822	0.816009	−0.272204	0.555728	−0.400228	
L_5	−0.003121	0.231629	−0.624239	0.129081	0.576830	−0.550403
R_5	−0.542431	−0.244341	0.024786	0.284971	−0.022551	−0.671125

Final combination:

$$\widetilde{T}_{24} = I + 0.984746\,B + 0.861176\,P_1 - 0.786548\,P_2 - 0.012111\,P_3 + 0.400063\,P_4 + P_5.$$

The prefix is exact to double precision: $\max_{0\le j\le 23} |[\widetilde{T}_{24}]_j - 1| < 5\times10^{-13}$ after local refinement. The refined spillover begins at

$$0.644829792\,B^{24} + 0.213865234\,B^{25} + 0.048460254\,B^{26} + 0.187929023\,B^{27} + \cdots + 0.056944198\,B^{32},$$

giving error-map leading coefficient $c = 1 - [\widetilde{T}_{24}]_{24} \approx 0.35517$.

Appendix 4. Certified Approximate Kernels

The rounded circuit tables in Appendix Appendix 2 are for display only. The authoritative data are the machine-readable supplementary files `certified_coeffs_radix15.json` and `certified_coeffs_radix24.json`, generated by `certify_kernel.py`.

Certification workflow. For each radix-m ansatz we form the prefix map

$$F_m(\theta) = ([f_\theta]_0 - 1, \ldots, [f_\theta]_{m-1} - 1).$$

At a numerically discovered seed θ_0, we compute the full Jacobian $DF_m(\theta_0)$ and select a well-conditioned square active subsystem. Holding the remaining variables fixed, we apply a Krawczyk inclusion test on a small box in the active coordinates. This certifies the existence of an exact

root of the prefix equations near the discovered seed. Evaluating the polynomial and error map over that certified box then yields interval enclosures for the spillover and error-map coefficients.

Certified kernels. For radix 15, the active-box radius is below 4.5×10^{-14} and the certified error map satisfies

$$E_{15}(z) = e_{15}z^{15} + e_{16}z^{16} + e_{17}z^{17},$$

with intervals centered at

$$e_{15} = 0.8147910024806893,$$
$$e_{16} = -0.2731510277211418,$$
$$e_{17} = 0.4583600252404543.$$

For radix 24, the active-box radius is below 1.5×10^{-10} and the certified error map coefficients are centered at

$$e_{24} = 0.3551702076534789, \quad e_{25} = 0.4309645579871833,$$
$$e_{26} = 0.1654049800790244, \quad e_{27} = -0.1394687687945098,$$
$$e_{28} = -0.2416800071277084, \quad e_{29} = -0.1015904946346608,$$
$$e_{30} = 0.1098721022155451, \quad e_{31} = 0.2068198204262950,$$
$$e_{32} = 0.1575634044917271, \quad e_{33} = 0.0569441977034457.$$

The exact certified intervals, including radii, are stored in the supplementary JSON files.

Certified radii. The coefficient-surrogate radii are certified at

$$\rho_{\eta,15} = 0.970897959\ldots, \qquad \rho_{\eta,24} = 0.974559738\ldots.$$

The direct circle-maximization test with coefficient intervals certifies

$$\rho^{\mathrm{cert}}_{q,15} = 0.970699488\ldots, \qquad \rho^{\mathrm{cert}}_{q,24} = 0.984176086\ldots.$$

For comparison, midpoint circle maximization gives the estimated values $0.970897959\ldots$ and $0.984432110\ldots$.

Appendix 5. Local Structure of the Prefix-Exact Set

The certification step yields concrete base points at which the prefix maps have full row rank. This supports a local structural theorem for the approximate-kernel ansätze.

Proposition Appendix 5.1 (Local structure near the certified kernels)**.** *Let F_m denote the prefix map for one of the certified bilinear ansätze, with parameter count p. At the certified radix-15 and radix-24 base points, the Jacobian DF_m has full row rank m. Consequently, the prefix-exact set $F_m(\theta) = 0$ is locally a smooth manifold of dimension $p - m$: 13 for radix 15 and 17 for radix 24. If the leading spillover coefficient is appended as an additional constraint, the local dimension drops by one.*

Proof. The rank computations are reported in `jacobian_rank_re` and are reproduced by `certify_kernel.py`. Full row rank of DF_m invokes the implicit function theorem. Appending one independent spillover constraint increases the rank by one for both certified kernels, so the local dimension decreases by one. □

Remark Appendix 5.2. This is a local statement only. It does not prove the global nonexistence of exact radix-15 or radix-24 kernels, nor does it exclude exact kernels in other ansätze.

Appendix 6. Companion Software and Reproducibility

The manuscript now depends on a reproducible software workflow:

- `compute_qrho.py` refreshes the certified radius artifacts;
- `certify_kernel.py` generates the certified coefficient files, the certified radii, and the Jacobian-rank report;
- `reproduce_theory_figures.py` regenerates the theory-driven figures and supporting CSV tables;
- `reproduce_all.py` reruns the certification scripts and then rebuilds the manuscript PDF;
- `reproducibility_manifest.txt` records the mapping from each generated figure or table to the script that produced it.

Machine-readable artifacts. The supplement contains:

- `certified_coeffs_radix15.json` and `certified_coeffs_`
- `qrho_radii.csv` and `jacobian_rank_report.csv`;
- `theory_validation.csv` and `failure_side.csv`;
- `reproducibility_manifest.txt`.

Companion repository. The public companion repository is planned at the project GitHub repository. It is intended to contain the manuscript sources, certified data files, tests, figure-reproduction scripts, and a copy of the current arXiv PDF. The companion tutorial site is planned at the tutorial site. The refreshed arXiv release and the repository release will be synchronized by tag so the posted source bundle and the public code snapshot refer to the same artifacts.